


Deadlocks

Thomas Plagemann

With slides from C. Griwodz, K. Li,
A. Tanenbaum and M. van Steen


Preempting Scheduler Activations

- Scheduler activations are completely preemptable



Preempting Scheduler Activations

- Maintaining the run queue needs to be a protected critical section
- Let's use spin locks for protection



Resources

- Examples of computer resources
 - CPU
 - Memory
 - Disk drive
 - Tape drives
 - Printers
 - Plotter
 - Loudspeaker

Resources

- Processes
 - Need access to resources in reasonable order
- Typical way to use a resource
 - Request
 - Use
 - Release
- Suppose a process holds resource A and requests resource B
 - At same time another process holds B and requests A
 - Both are blocked and remain so

Resources



- Active resource
 - Provides a service
 - E.g. CPU, network adaptor
- Passive resource
 - System capabilities that are required by active resources
 - E.g. memory, network bandwidth
- Exclusive resource
 - Only one process at a time can use it
 - E.g. loudspeaker, processor
- Shared resource
 - Can be used by multiple processes
 - E.g. memory, bandwidth

Resources

- Single resource of each type
 - Exists only once in the system
 - E.g. loudspeaker
- Multiple resources of each type
 - Several instances of resource in the system
 - E.g. processor in a multiprocessor system
- Preemptable resource
 - Resource that can be taken away from a process
 - E.g. CPU can be taken away from processes in user space
- Non-preemptable resource
 - Taking it away will cause processes to fail
 - E.g. Disk, files
- What would you need to model resources of overlapping capabilities, e.g. color and B/W printer?

Resources

- Process must wait if request is denied
 - Requesting process may be blocked
 - May fail with error code
- Deadlocks
 - Can occur only if processes are granted exclusive access to resources





Deadlocks

- Formal definition :

*A set of processes is deadlocked
if each process in the set is waiting for an event
that only another process in the set can cause*

- Usually the *event* is release of a currently held resource
- None of the processes can ...
 - Run
 - Release resources
 - Be awakened



Four Conditions for Deadlock

1. Mutual exclusion condition

- Each resource assigned to 1 process or is available

2. Hold and wait condition

- Process holding resources can request additional

3. No preemption condition


- Previously granted resources cannot forcibly taken away

4. Circular wait condition

- Must be a circular chain of 2 or more processes
- Each is waiting for resource held by next member of the chain

Deadlock Modeling

- Modeled with directed graphs




- Resource R assigned to process A
- Process B is requesting/waiting for resource S
- Process C and D are in deadlock over resources T and U

Deadlock Example

- A utility program
 - Copies a file from a tape to disk
 - Prints the file to a printer
- Resources
 - Tape
 - Disk
 - Printer

- A deadlock



Deadlock Modeling

■ How deadlock occurs


A
Requests R
Requests S
Releases S
Releases R

B
Requests S
Requests T
Releases T
Releases S

C
Requests T
Requests R
Releases R
Releases T

Processes

Resources



A requests R
B requests S
C requests T
A requests S
B requests T
C requests R

Deadlock Modeling

■ How deadlock can be avoided


A
Requests R
Requests S
Releases S
Releases R

B
Requests S
Requests T
Releases T
Releases S

C
Requests T
Requests R
Releases R
Releases T

Processes

Resources



A requests R
C requests T
A requests S
B requests S
B requests T
C requests R
A releases S
A releases R
C releases R
C releases T



Deadlocks: Strategies

- Ignore the problem
 - It is user's fault
- Detection and recovery
 - Fix the problem afterwards
- Dynamic avoidance
 - Careful allocation
- Prevention
 - Negate one of the four conditions

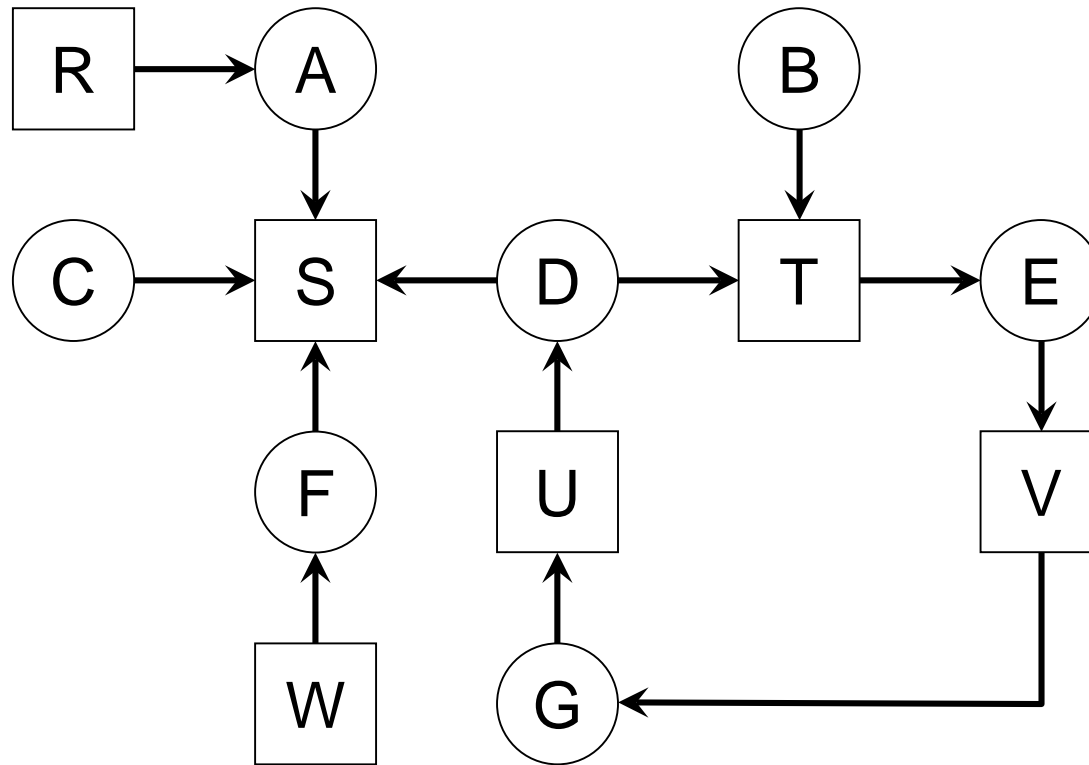
The Ostrich Algorithm

- Pretend there is no problem
- Reasonable if
 - Deadlocks occur very rarely
 - Cost of prevention is high
- UNIX and Windows take this approach
- It is a trade-off between
 - Convenience
 - Correctness



Deadlock Detection and Recovery

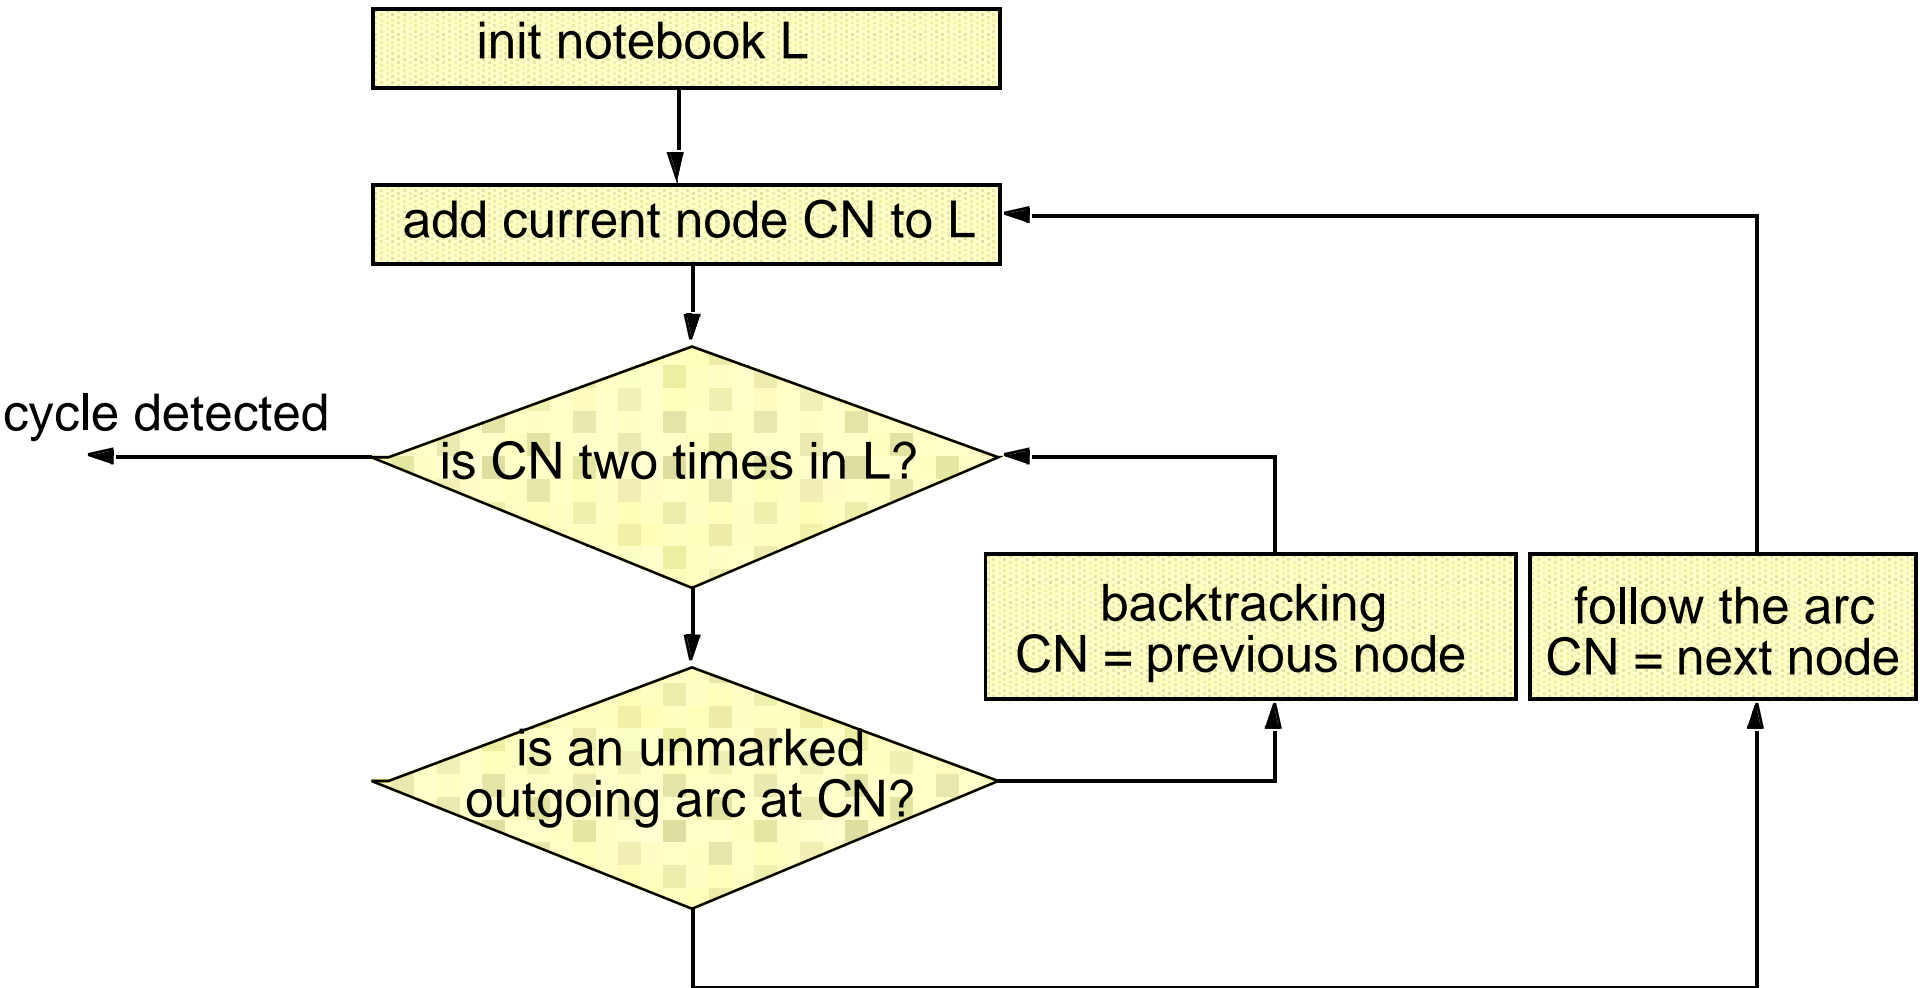
One Resource of Each Type



- A cycle can be found within the graph, denoting deadlock

Deadlock Detection and Recovery

One Resource of Each Type



Deadlock Detection and Recovery

Multiple Resources of Each Type

Existing resources
 $(E_1, E_2, E_3, \dots, E_m)$

Current allocation matrix

$$\begin{bmatrix} C_{11} & C_{12} & C_{13} & \dots & C_{1m} \\ C_{21} & C_{22} & C_{23} & \dots & C_{2m} \\ \dots & \dots & \dots & \dots & \dots \\ C_{n1} & C_{n2} & C_{n3} & \dots & C_{nm} \end{bmatrix}$$

Available resources
 $(A_1, A_2, A_3, \dots, A_m)$

Request matrix

$$\begin{bmatrix} R_{11} & R_{12} & R_{13} & \dots & R_{1m} \\ R_{21} & R_{22} & R_{23} & \dots & R_{2m} \\ \dots & \dots & \dots & \dots & \dots \\ R_{n1} & R_{n2} & R_{n3} & \dots & R_{nm} \end{bmatrix}$$

Deadlock Detection and Recovery

Multiple Resources of Each Type

$$E = \begin{pmatrix} & \begin{matrix} \text{Tape drivers} \\ \text{Plotters} \\ \text{Scanners} \\ \text{CD-Roms} \end{matrix} \\ \begin{matrix} 4 \\ 2 \\ 3 \\ 1 \end{matrix} & \end{pmatrix}$$

$$A = \begin{pmatrix} & \begin{matrix} \text{Tape drivers} \\ \text{Plotters} \\ \text{Scanners} \\ \text{CD-Roms} \end{matrix} \\ \begin{matrix} 2 \\ 1 \\ 0 \\ 0 \end{matrix} & \end{pmatrix}$$

Current allocation matrix

$$C = \begin{bmatrix} 0 & 0 & 1 & 0 \\ 2 & 0 & 0 & 1 \\ 0 & 1 & 2 & 0 \end{bmatrix}$$

Request matrix


$$R = \begin{bmatrix} 2 & 0 & 0 & 1 \\ 1 & 0 & 1 & 0 \\ 2 & 1 & 0 & 0 \end{bmatrix}$$

Deadlock Detection and Recovery

Recovery

- Recovery through preemption
 - Take a resource from some other process
 - Depends on nature of the resource
- Recovery through rollback
 - Checkpoint a process periodically
 - Use this saved state
 - Restart the process if it is found deadlocked
- Recovery through killing processes
 - Crudest but simplest way to break a deadlock
 - Kill one of the processes in the deadlock cycle
 - The other processes get its resources
 - Choose process that can be rerun from the beginning

Two process resource trajectories

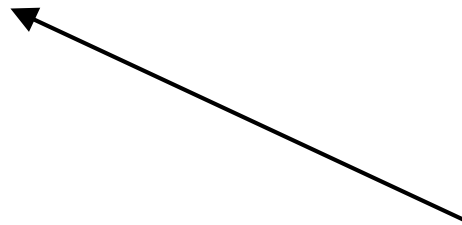


Deadlock Avoidance

Safe and Unsafe States

has max			has max			has max			has max			has max		
A	3	9	A	3	9	A	3	9	A	3	9	A	3	9
B	2	4	B	4	4	B	0		B	0		B	0	
C	2	7	C	2	7	C	2	7	C	7	7	C	0	
Free: 3			Free: 1			Free: 5			Free: 0			Free: 7		

state is safe

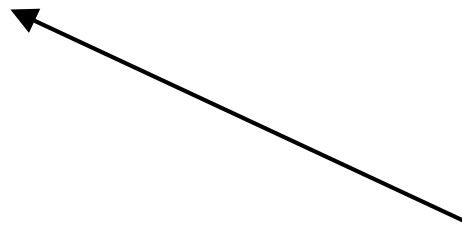


Deadlock Avoidance

Safe and Unsafe States

has max			has max			has max			has max		
A	3	9	A	4	9	A	4	9	A	3	9
B	2	4	B	2	4	B	4	4	B	0	
C	2	7	C	2	7	C	2	7	C	2	7
Free: 3			Free: 2			Free: 0			Free: 4		

state is safe



Deadlock Avoidance

Banker's Algorithm for a Single Resource

- Each process has a credit
 - System knows how many resources a process requests *at most* before releasing resources
- Total resources may not satisfy all credits
- Keep track of resources assigned and needed
- Check on each allocation whether it is safe
 - Safe: there exists a sequence of other states that all processes can terminate correctly

Deadlock Avoidance

Banker's Algorithm for a Single Resource

Resource allocation state

has max

A	0	6
B	0	5
C	0	4
D	0	7

Free: 10

has max

A	1	6
B	1	5
C	2	4
D	4	7

Free: 2

has max

A	1	6
B	2	5
C	2	4
D	4	7

Free: 1

Deadlock Detection and Recovery

Banker's Algorithm for Multiple Resources

Assigned resources

A	3	0	1	1
B	0	1	0	0
C	1	1	1	0
D	1	1	0	1
E	0	0	0	0

Resources still needed

A	1	1	0	0
B	0	1	1	2
C	3	1	0	0
D	0	0	1	0
E	2	1	1	0

$$\begin{array}{l}
 \text{Tape drives} \\
 \text{Plotters} \\
 \text{Scanners} \\
 \text{CD-Roms}
 \end{array}$$

$$\begin{array}{l}
 E = (\begin{array}{|c|c|c|c|} \hline 6 & 3 & 4 & 2 \\ \hline \end{array}) \\
 P = (\begin{array}{|c|c|c|c|} \hline 5 & 3 & 2 & 2 \\ \hline \end{array}) \\
 A = (\begin{array}{|c|c|c|c|} \hline 1 & 0 & 2 & 0 \\ \hline \end{array})
 \end{array}$$

An example for the deadlock detection algorithm

Deadlock Detection and Recovery

Banker's Algorithm for Multiple Resources

Assigned resources

	Tape drives	Plotters	Scanners	CD-Roms
A	3	0	1	1
B	0	1	0	0
C	1	1	1	0
D	0	0	0	0
E	0	0	0	0

Resources still needed

	Tape drives	Plotters	Scanners	CD-Roms
A	1	1	0	0
B	0	1	1	2
C	3	1	0	0
D	-	-	-	-
E	2	1	1	0

$$\begin{aligned}
 E &= \begin{pmatrix} 6 & 3 & 4 & 2 \end{pmatrix} \\
 P &= \begin{pmatrix} 4 & 2 & 2 & 1 \end{pmatrix} \\
 A &= \begin{pmatrix} 2 & 1 & 2 & 1 \end{pmatrix}
 \end{aligned}$$

An example for the deadlock detection algorithm

Deadlock Detection and Recovery

Banker's Algorithm for Multiple Resources

Assigned resources

A	0	0	0	0
B	0	1	0	0
C	1	1	1	0
D	0	0	0	0
E	0	0	0	0

Resources still needed

A	-	-	-	-
B	0	1	1	2
C	3	1	0	0
D	-	-	-	-
E	2	1	1	0

$$\begin{array}{l} \text{Tape drives} \\ \text{Plotters} \\ \text{Scanners} \\ \text{CD-Roms} \end{array}$$
$$E = \begin{pmatrix} 6 & 3 & 4 & 2 \end{pmatrix}$$
$$P = \begin{pmatrix} 1 & 2 & 1 & 0 \end{pmatrix}$$
$$A = \begin{pmatrix} 5 & 1 & 3 & 2 \end{pmatrix}$$

An example for the deadlock detection algorithm

Deadlock Detection and Recovery

Banker's Algorithm for Multiple Resources

Assigned resources

A	0	0	0	0
B	0	0	0	0
C	1	1	1	0
D	0	0	0	0
E	0	0	0	0

Resources still needed

A	-	-	-	-
B	-	-	-	-
C	3	1	0	0
D	-	-	-	-
E	2	1	1	0

$$\begin{array}{l}
 \text{Tape drives} \\
 \text{Plotters} \\
 \text{Scanners} \\
 \text{CD-Roms}
 \end{array}$$

$$\begin{array}{l}
 E = (\begin{array}{|c|c|c|c|} \hline 6 & 3 & 4 & 2 \\ \hline \end{array}) \\
 P = (\begin{array}{|c|c|c|c|} \hline 1 & 1 & 1 & 0 \\ \hline \end{array}) \\
 A = (\begin{array}{|c|c|c|c|} \hline 5 & 2 & 3 & 2 \\ \hline \end{array})
 \end{array}$$

An example for the deadlock detection algorithm

Deadlock Detection and Recovery

Banker's Algorithm for Multiple Resources

Assigned resources

A	0	0	0	0
B	0	0	0	0
C	0	0	0	0
D	0	0	0	0
E	0	0	0	0

Resources still needed

A	-	-	-	-
B	-	-	-	-
C	-	-	-	-
D	-	-	-	-
E	2	1	1	0

$$\begin{array}{l} \text{Tape drives} \\ \text{Plotters} \\ \text{Scanners} \\ \text{CD-Roms} \end{array}$$
$$E = \begin{pmatrix} 6 & 3 & 4 & 2 \end{pmatrix}$$
$$P = \begin{pmatrix} 0 & 0 & 0 & 0 \end{pmatrix}$$
$$A = \begin{pmatrix} 6 & 3 & 4 & 2 \end{pmatrix}$$

An example for the deadlock detection algorithm

Deadlock Detection and Recovery

Banker's Algorithm for Multiple Resources

SAFE

Assigned resources

	Tape drives	Plotters	Scanners	CD-Roms
A	3	0	1	1
B	0	1	0	0
C	1	1	1	0
D	1	1	0	1
E	0	0	0	0

Resources still needed

	Tape drives	Plotters	Scanners	CD-Roms
A	1	1	0	0
B	0	1	1	2
C	3	1	0	0
D	0	0	1	0
E	2	1	1	0

$$\begin{aligned}
 E &= (\begin{array}{|c|c|c|c|} \hline 6 & 3 & 4 & 2 \\ \hline \end{array}) \\
 P &= (\begin{array}{|c|c|c|c|} \hline 5 & 3 & 2 & 2 \\ \hline \end{array}) \\
 A &= (\begin{array}{|c|c|c|c|} \hline 1 & 0 & 2 & 0 \\ \hline \end{array})
 \end{aligned}$$

An example for the deadlock detection algorithm

Deadlock Avoidance

Practical Avoidance

- Two Phase Locking

- Phase I

- Process tries to lock all resources it needs, one at a time
 - If needed resources found locked, start over
 - (no real work done in phase one)

- Phase II

- Run
 - Releasing locks

- Note similarity to requesting all resources at once
- Algorithm works where programmer can arrange

Deadlock Prevention

R: Conditions for Deadlock

1. Mutual exclusion condition

- Each resource assigned to 1 process or is available

2. Hold and wait condition

- Process holding resources can request additional

3. No preemption condition

- Previously granted resources cannot forcibly taken away

4. Circular wait condition

- Must be a circular chain of 2 or more processes
- Each is waiting for resource held by next member of the chain

Deadlock Prevention

Mutual Exclusion Condition

- Some resources are not sharable
 - Printer, tape, etc
- Some resources can be made sharable
- Some resources can be made virtual
 - Spooling - Printer
 - Does spooling apply to all non-sharable resources?
 - Mixing - Soundcard
- Principle:
 - Avoid assigning resource when not absolutely necessary
 - A few processes as possible actually claim the resource

Deadlock Prevention


Hold and Wait Condition

- Require processes to request resources before starting
 - A process never has to wait for what it needs
 - Telephone companies do this
- Problems
 - May not know required resources at start of run
 - Also ties up resources other processes could be using
- Variation:
 - Process must give up all resources
 - Then request all immediately needed

Deadlock Prevention

No Preemption Condition


- This is not a viable option
- Consider a process given the printer
 - Halfway through its job
 - No forcibly take away printer
 - !!??



Deadlock Prevention

Circular Wait Condition

1. CD Rom drive
2. Tape drive
3. Plotter
4. Scanner
5. Imagesetter



- Normally ordered resources
- A resource graph

Deadlock Prevention

Circular Wait Condition

- Impose an order of requests for all resources
- Method
 - Assign a unique id to each resource
 - All resource requests must be in an ascending order of the ids
 - Release resources in a descending order
- Can you prove this method has no circular wait?
- Is this generally feasible?

Deadlock Prevention Overview


Condition	Approach
Mutual exclusion	Spool everything
Hold and wait	Request all resource initially
No preemption	Take resources away
Circular wait	Order resources numerically

Non-resource Deadlocks

- Possible for two processes to deadlock
 - Each is waiting for the other to do some task
- Can happen with semaphores
 - Each process required to do a *down()* on two semaphores (*mutex* and another)
 - If done in wrong order, deadlock results

Preempting Scheduler Activations

- So how do they handle this deadlock?



Preempting Scheduler Activations

- Detection and recovery (like in Mach)
- Basic idea:
 - Upcall handler checks first the state of each interrupted thread
 - If it is in a critical section allow it to finish this
- Implementation:
 - Protect critical sections with spin locks
 - `Acquire_spin_lock()` increments a counter in the thread's descriptor
 - Upcall handler checks spin lock count of all interrupted threads
 - If there are threads that hold spin locks set flag
 - Switch to the context of these threads
 - Afterwards switch back to upcall handler

Summary

- Resource
- Introduction to deadlocks
- Strategies
 - Ostrich algorithm
 - Deadlock detection and recovery
 - Deadlock avoidance
 - Deadlock prevention
- Non-resource deadlocks

Other Presentations

- Singh – Princeton

- <http://www.cs.princeton.edu/courses/archive/fall18/cos318/lectures/9.Deadlock.pdf>