

# Operating Systems Structure and Processes

Otto J. Anshus  
University of Tromsø/Oslo

# The Architecture of an OS

- Monolithic
- Layered
- Virtual Machine, Library, Exokernel
- Micro kernel and Client/Server
- Hybrids


# Goals of the architecture

- OS as Resource Manager
- OS as Virtual Machine (abstractions)
- Efficiency, flexibility, size, security, ... as discussed earlier


# Operating System Use of Processes: *Where is the OS executing?*

*Illustrations: [Stallings:  
Operating System:  
Internals and Design](#)*

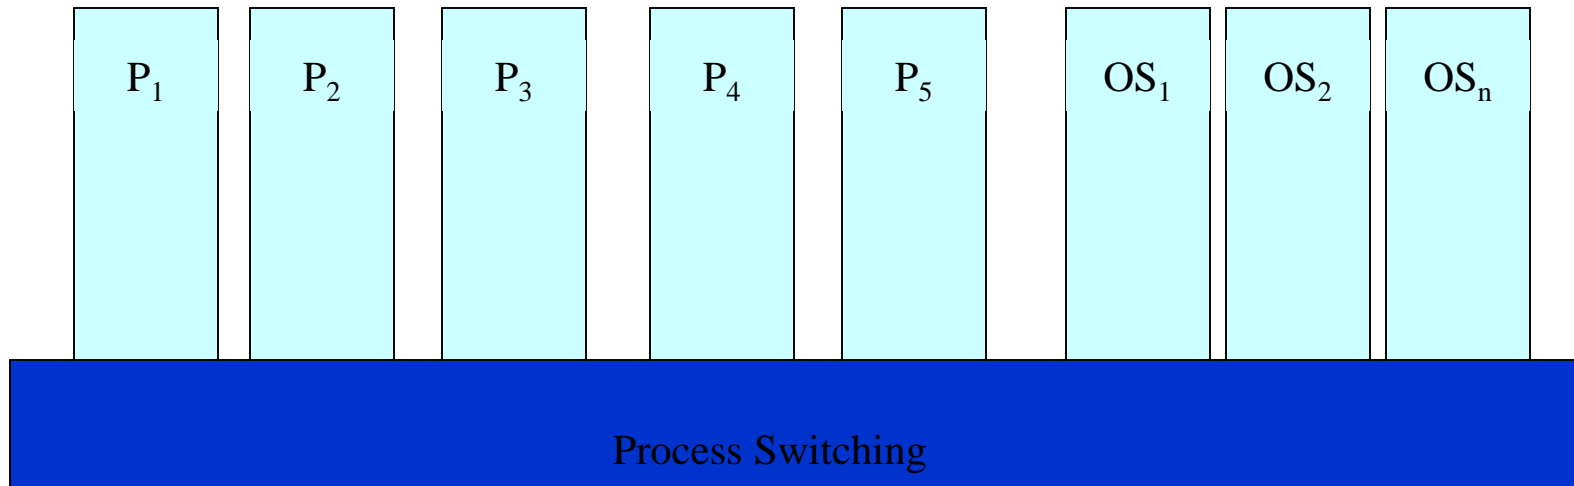
## Separate Kernel

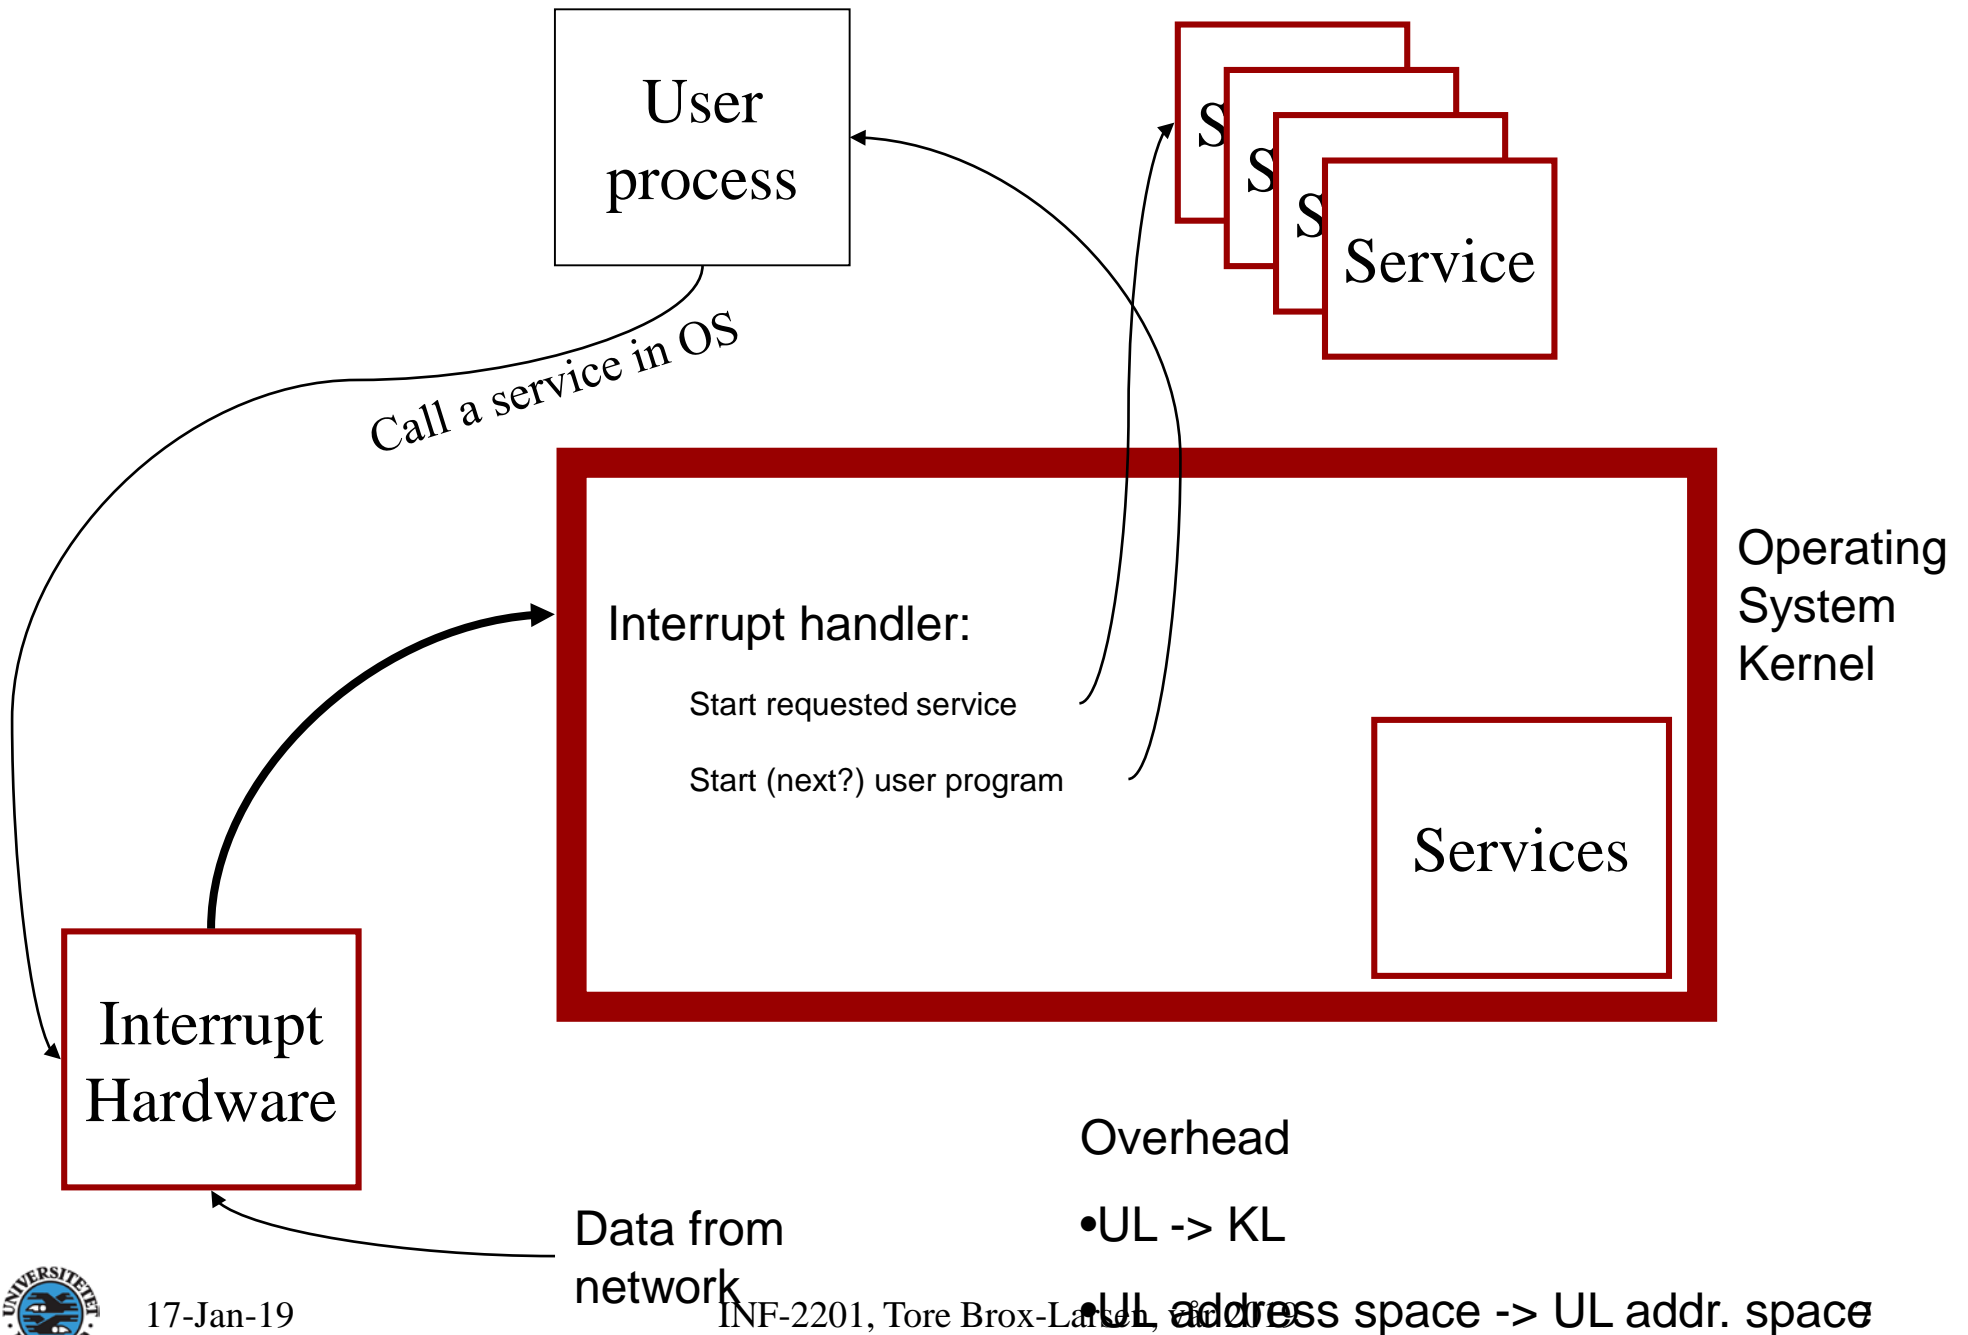


# OS-Functions Executing within Processes




# OS-Functions Executing in Separate Processes





# Monolithic

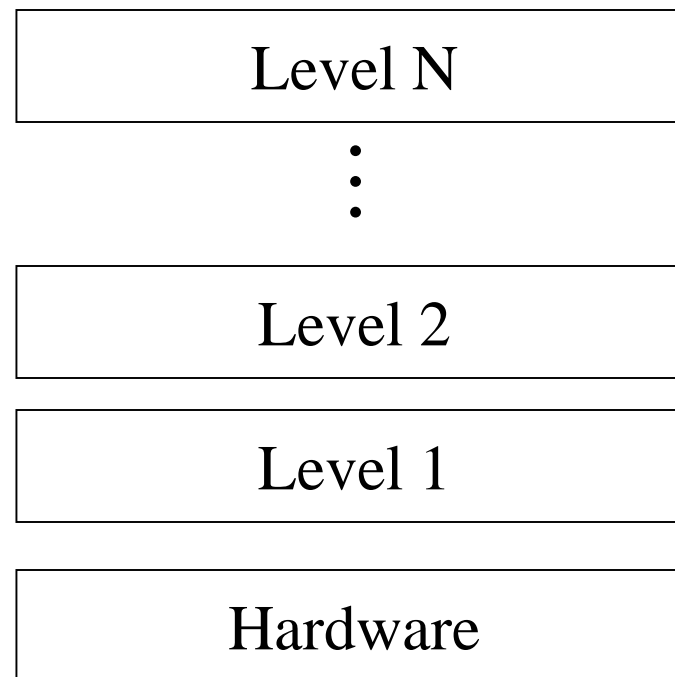
- All kernel routines are together
- A system call interface
- Examples:
  - Linux
  - Most Unix OS
  - NT (hybrid)





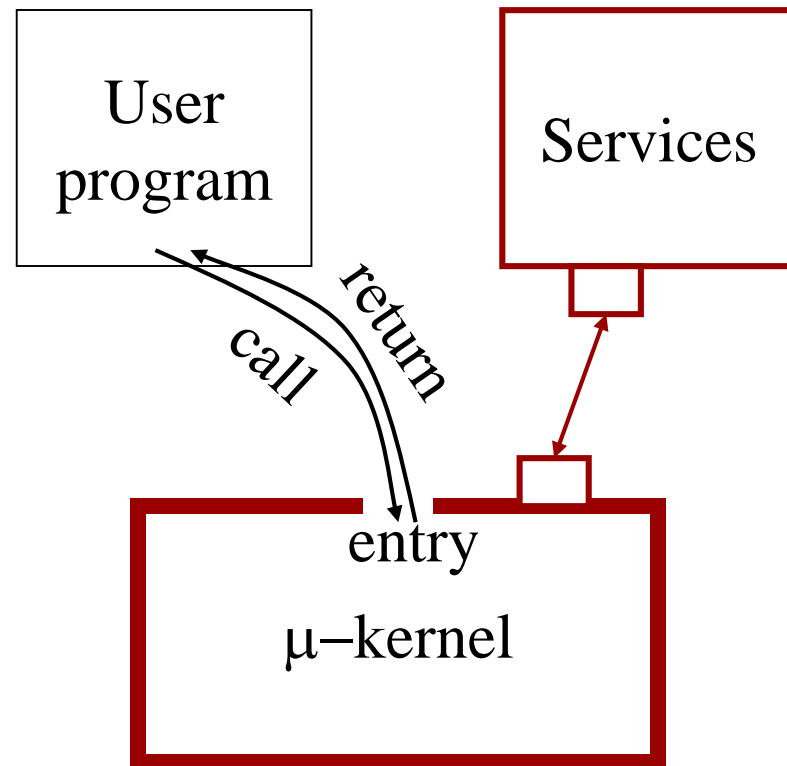
# Layered Structure

- Hiding information at each layer
- Develop a layer at a time
- Examples
  - [THE](#) (6 layers, semaphores, [Dijkstra](#) 1968)
  - [MS-DOS](#) (4 layers)



# Microkernel and Client/Server

- Micro-kernel is “micro”
- Services are implemented as user level processes
- Micro-kernel get services on behalf of users by messaging with the service processes
- Example: [L4](#), [Nucleus](#), [Taos](#), [Mach](#), [Mach](#), [NT](#) (hybrid)



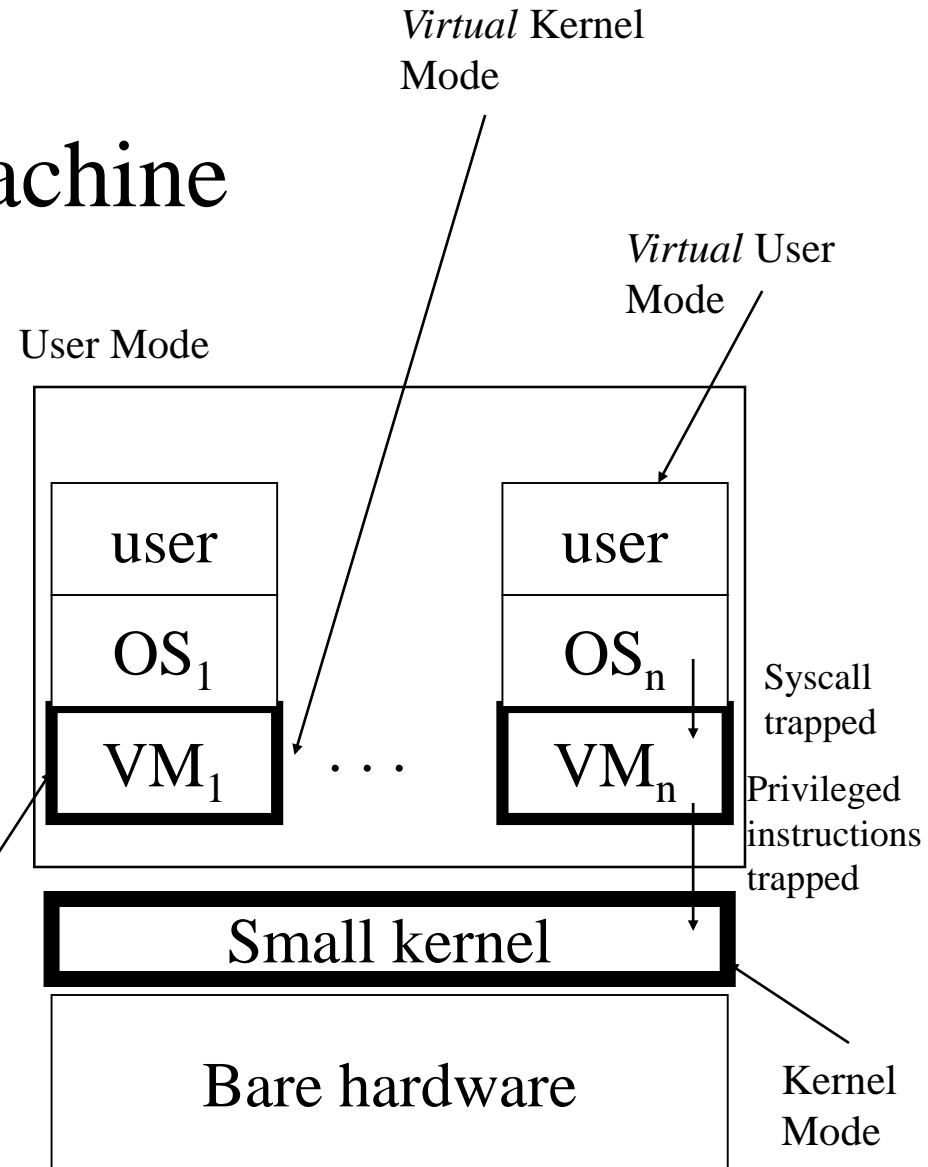
# Virtual Machine

"A running program is often referred to as a virtual machine - a machine that doesn't exist as a matter of actual physical reality. The virtual machine idea is itself one of the most elegant in the history of technology and is a crucial step in the evolution of ideas about software. To come up with it, scientists and technologists had to recognize that a computer running a program isn't merely a washer doing laundry. A washer is a washer whatever clothes you put inside, but when you put a new program in a computer, it becomes a new machine.... The virtual machine: A way of understanding software that frees us to think of software design as machine design."

From David Gelernter's "[Truth, Beauty, and the Virtual Machine](#)," Discover Magazine, September 1997, p. 72.

# Virtual Machine

- Virtual machine monitor
  - provide multiple virtual “real” hardware
  - run different OS codes
- Example
  - [IBM VM/370](#): Started in the 70’s. [Check out](#)
  - virtual 8086 mode
  - Java VM
  - VMware
  - [Exokernel](#)



Input/Output

Console

Output

Auxiliary  
Storage

Processing and  
Main Storage

Input


Input

Output



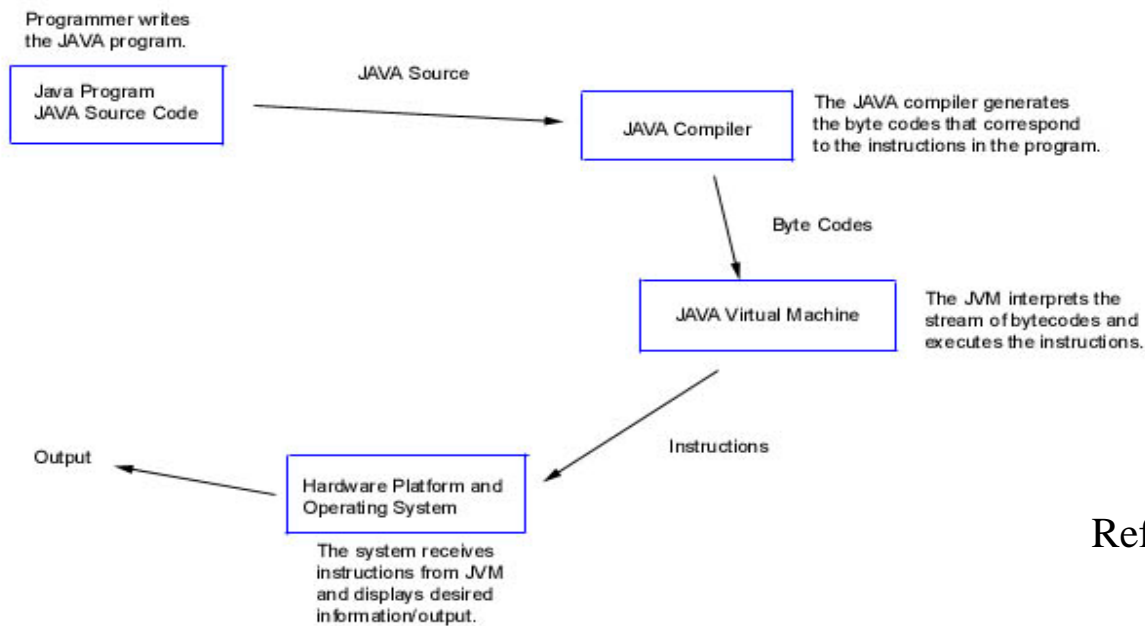
# Virtual 8086

## A NEW OLD IDEA: PENTIUM VIRTUAL 8086 MODE



- Virtual 8086 mode on the Pentium makes it possible to run old 16-bit DOS applications on a virtual machine


# Java VM



Ref. [Pascal P-Code](#)

Figure 1.1: Diagram of Java Program Execution

# Exokernel Architecture





# Hardware Support

- What is the minimal support?
  - 2 modes
  - Exception and interrupt trapping
- Can virtual machine be protected without such support?
  - Yes, emulation instead of executing on real machine

# Pro et Contra

Monolithic	Layered	VM	C/S	Micro kernel
<ul style="list-style-type: none"> <li>•Performance</li> </ul>	<ul style="list-style-type: none"> <li>•Clean, less bugs</li> <li>•Clear division of labour</li> </ul>	<ul style="list-style-type: none"> <li>•Many virtual computers with different OS'es</li> <li>•Test of new OS while production work continues</li> <li>•All in all: flexibility</li> </ul>	<ul style="list-style-type: none"> <li>•Clear division of labour</li> </ul>	<ul style="list-style-type: none"> <li>•More flexible</li> <li>•Small means less bugs+manageable</li> <li>•Distributed systems</li> <li>•Failure isolation of services at Kernel Level</li> </ul>
<ul style="list-style-type: none"> <li>•Less structured</li> </ul>	<ul style="list-style-type: none"> <li>•Are layers really separated?</li> <li>•Performance issues?</li> </ul>	<ul style="list-style-type: none"> <li>•Performance issues?</li> <li>•Complexity issues?</li> </ul>	<ul style="list-style-type: none"> <li>•Performance issues?</li> </ul>	<ul style="list-style-type: none"> <li>•Flexibility issues?</li> <li>•Performance issues?</li> </ul>

# “Truths” on Micro Kernel Flexibility and Performance

- A micro kernel restricts application level flexibility.
- Switching overhead kernel-user mode is inherently expensive.
- Switching address-spaces is costly.
- IPC is expensive.
- Micro kernel architectures lead to memory fragmentation.
- Kernel should be portable (on top of a small hardware-dependent layer).

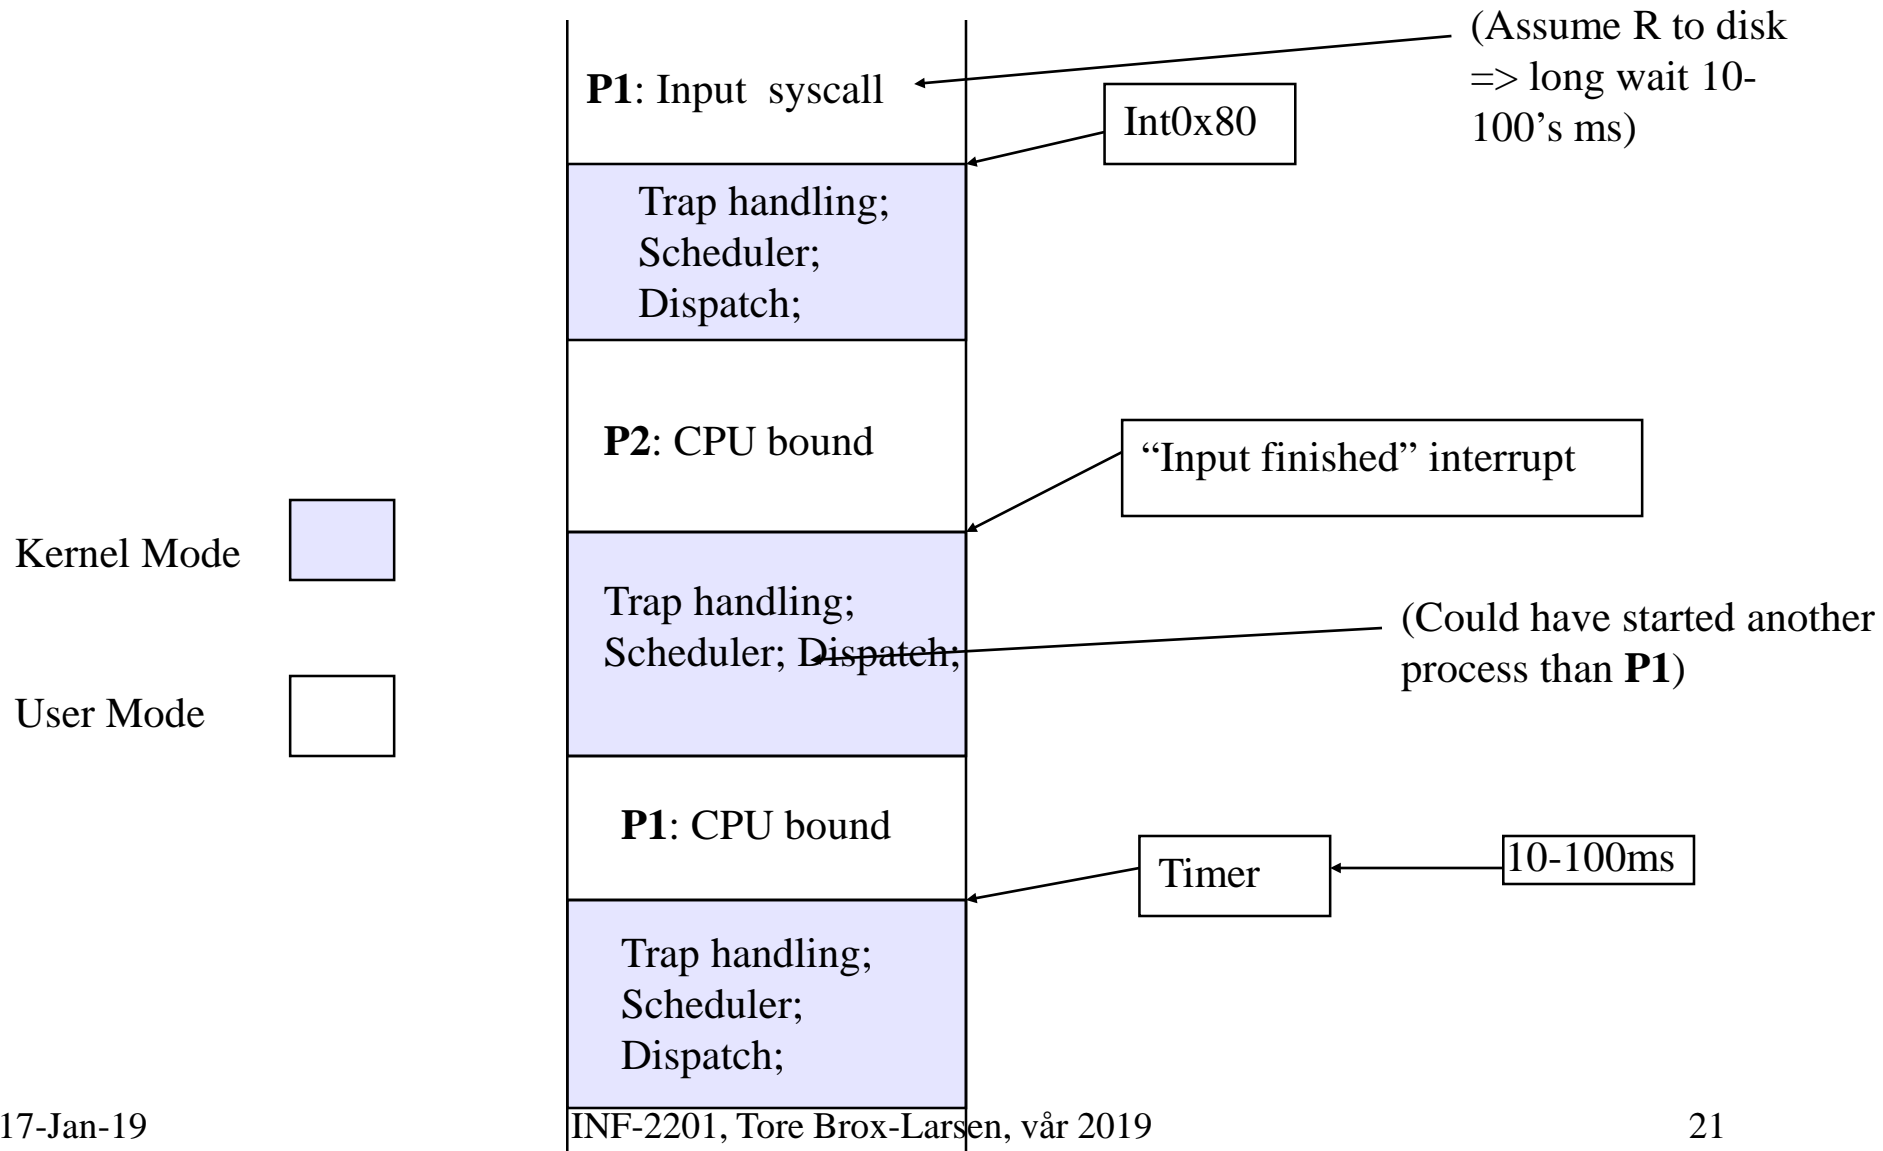
NO: Can be <50 cycles

NO: 6-20 microsec round-trip,  
53-500 cycles/IPC one way

# Concurrency and Process

- Problem to solve
  - A shared CPU, many I/O devices and lots of interrupts
  - Users feel they have machine to themselves
- Strategy
  - Decompose hard problems into simple ones
  - Deal with one at a time
  - Process is such a unit

# Flow of Execution



# Procedure, Co-routine, Thread, Process

- Procedure, Function, (Sub)Routine

- Call-execute all-return nesting

- Co-routine

- Call-resumes-return

← User level non preemptive “scheduler”  
in user code


- Thread (more later)

- Process


- Single threaded      ↓

- Multi threaded      ↓ ↓ ↓

# Procedure and Co-routine



“User Yield when finished”



“User Yield during execution to share CPU”

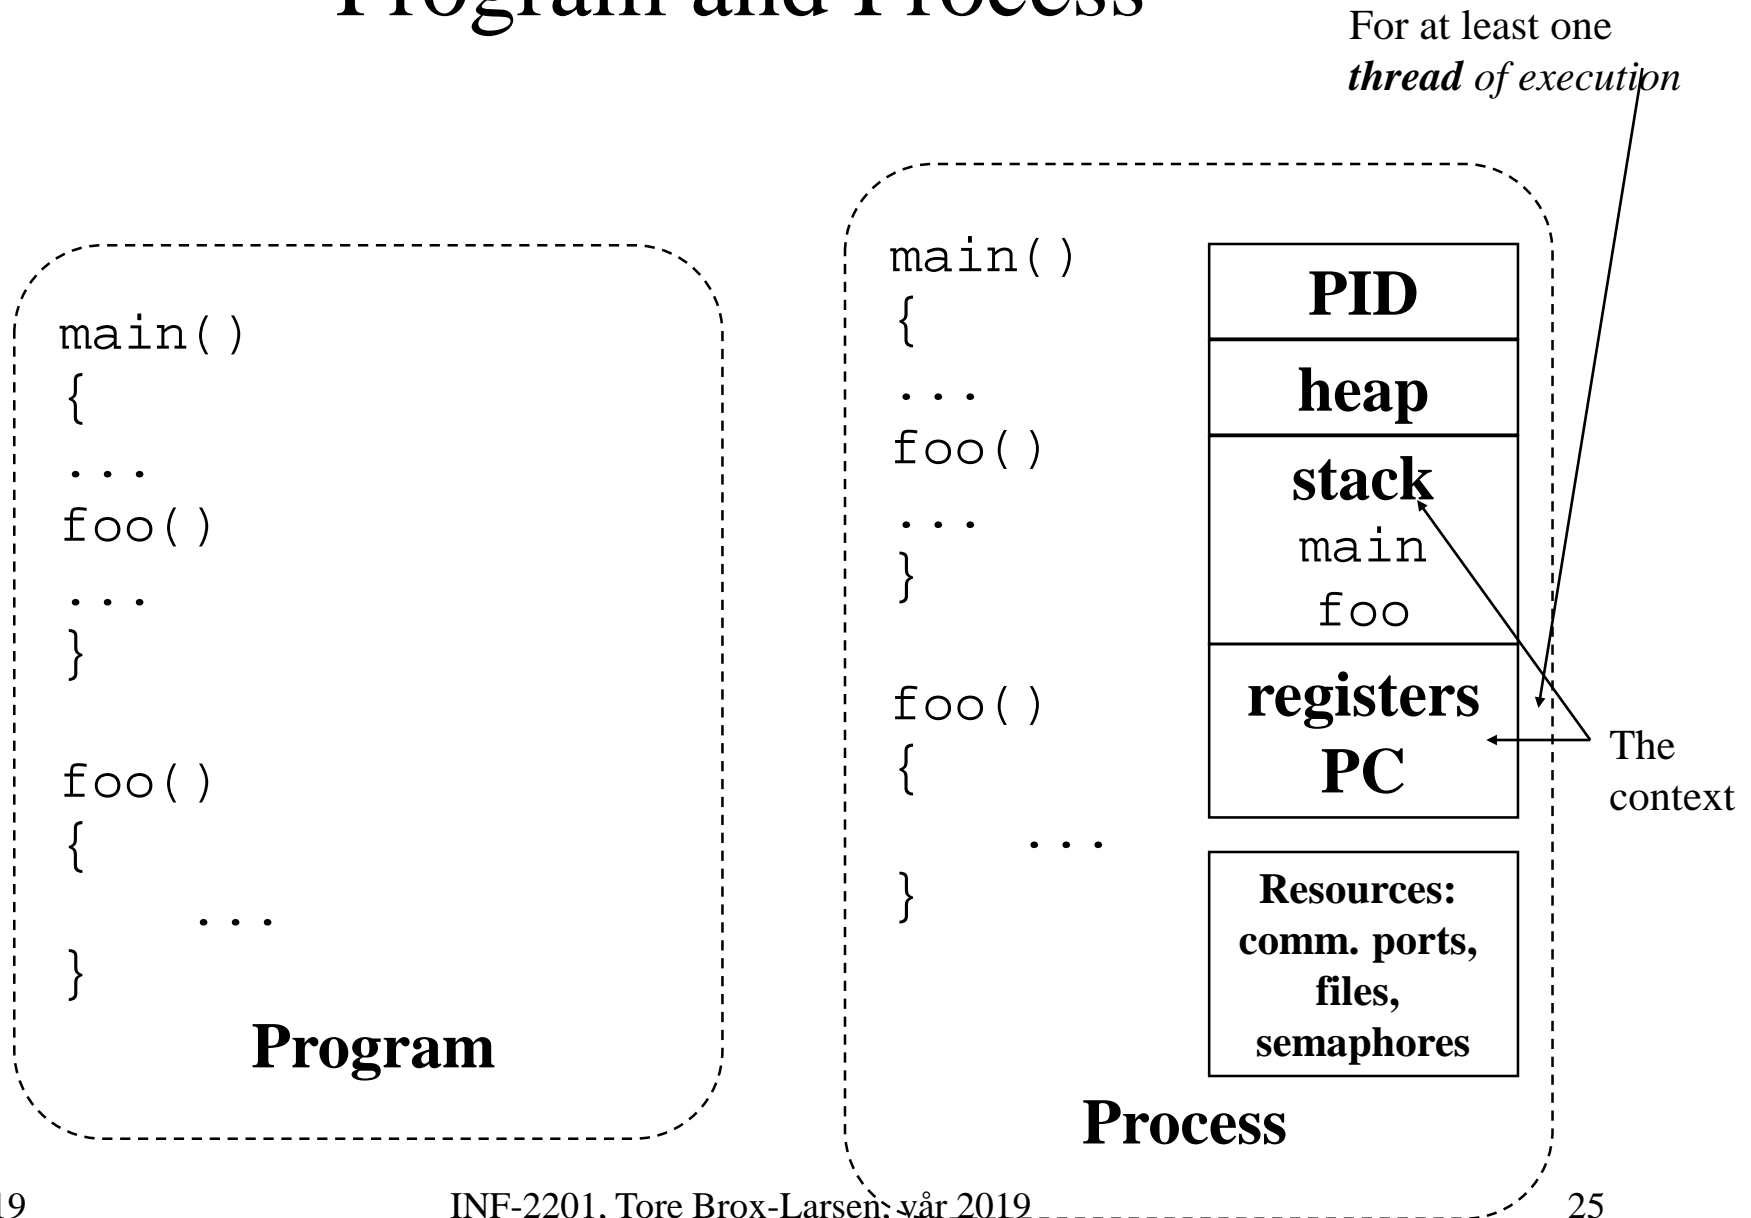
Never executed

# Process

- Sequential execution of operations
  - No concurrency inside a (**single** threaded) process
  - Everything happens sequentially
- Process state
  - Registers
  - Stack(s)
  - Main memory
  - Open files in UNIX
  - Communication ports
  - Other resources



# Program and Process

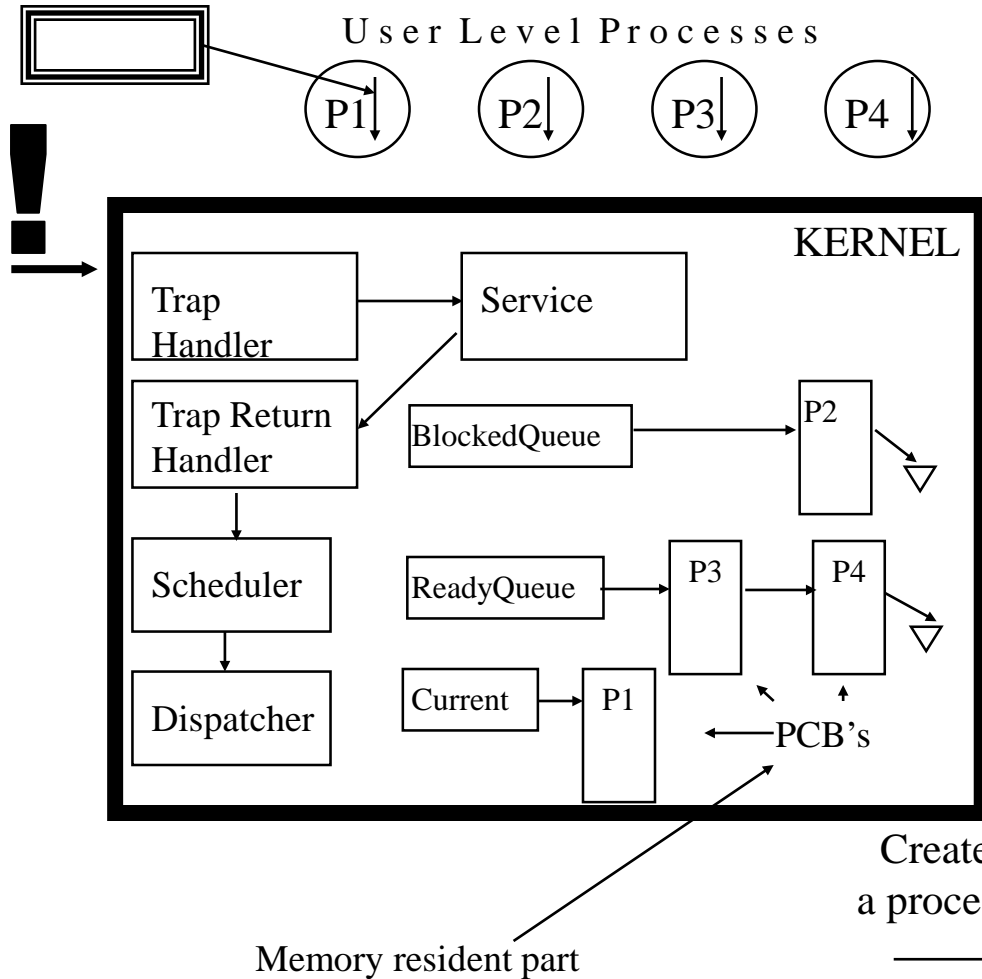


# Process vs. Program

- Process  $>$  program
  - Program is just part of process state
  - Example: many users can run the same program
- Process  $<$  program
  - A program can invoke more than one process
  - Example: Fork off processes to lookup webster


Instruction Pointer  
(program counter) in the  
EIP register

# Process State Transitions

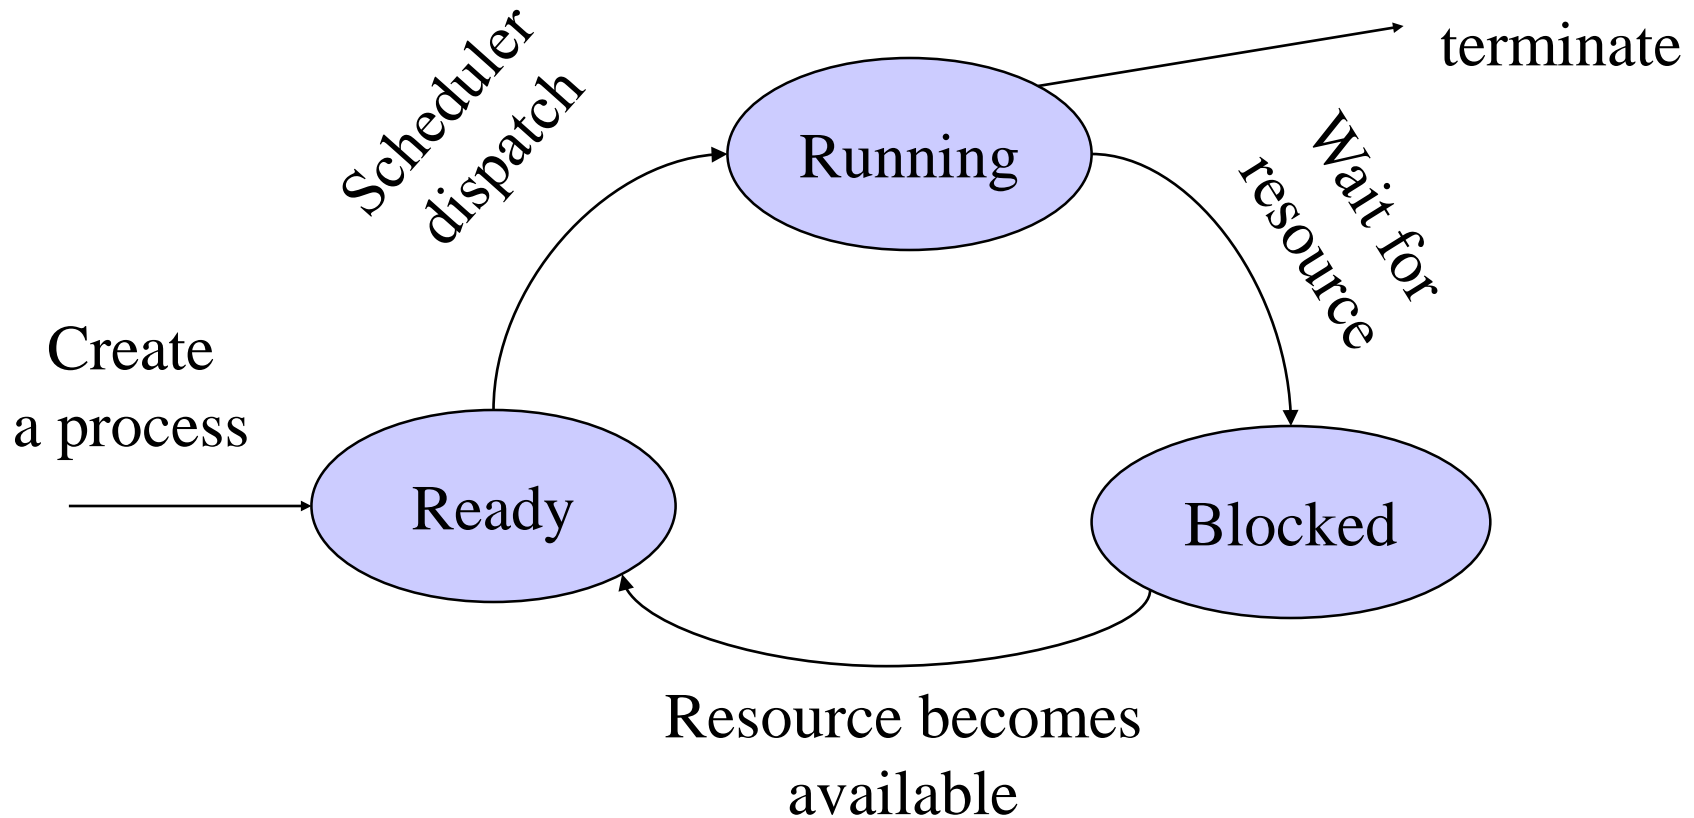


## MULTIPROGRAMMING

- Uniprocessor: *Interleaving* (“pseudoparallelism”)
- Multiprocessor: *Overlapping* (“true parallelism”)



# Process State Transition



# Process Control Block (Process Table)

- What
  - Process management info
    - State (ready, running, blocked)
    - Registers, PSW, parents, etc
  - Memory management info
    - Segments, page table, stats, etc
  - I/O and file management
    - Communication ports, directories, file descriptors, etc.

# Discussion: What needs to be saved and restored on a context switch?


- Volatile state
  - Program counter (Program Counter (PC) also called Instruction Pointer (Intel: EIP))
  - Processor status register
  - Other register contents
  - User and kernel stack pointers
  - A pointer to the address space in which the process runs
    - the process's page table directory

...and how?

- **Save**(volatile machine state, current process);
- **Load**(another process's saved volatile state);
- **Start**(new process);

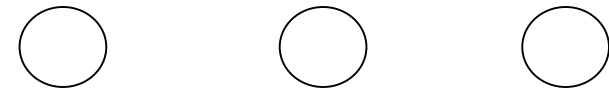
# Threads and Processes

## *Trad. Threads*

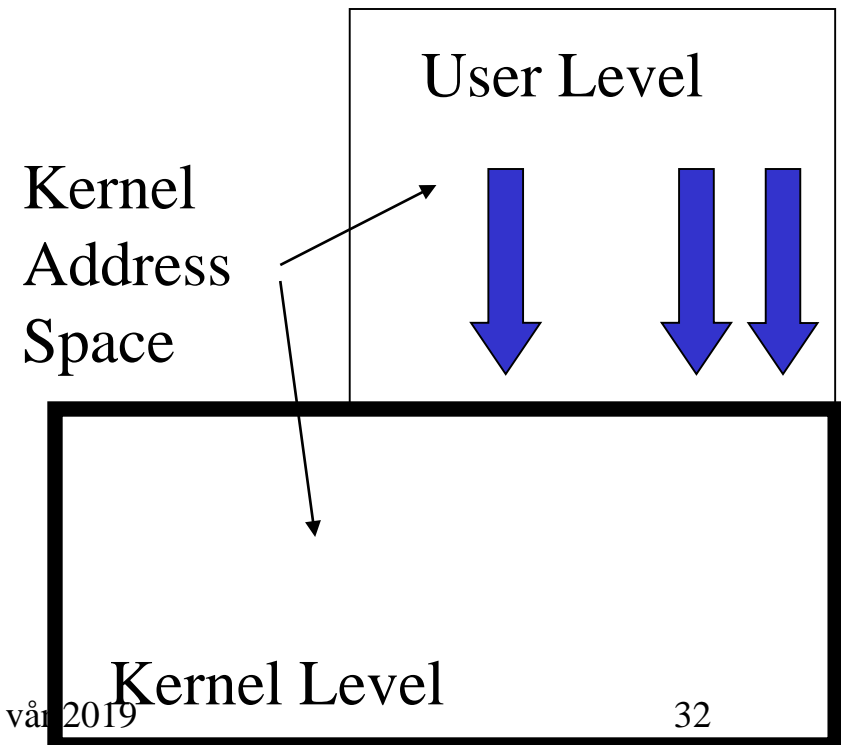


## *Project OpSys*

Processes in individual address spaces



Kernel threads





# Links

- Virtual machine

- [http://whatis.techtarget.com/definition/0,,sid9\\_gci213305,00.html](http://whatis.techtarget.com/definition/0,,sid9_gci213305,00.html)

- Exokernel

- <https://en.wikipedia.org/wiki/Exokernel>

- THE

- <http://www.cs.utexas.edu/users/EWD/ewd01xx/EWD196.PDF>

- L4

- <http://os.inf.tu-dresden.de/L4/>

- VM

- <http://www.vm.ibm.com/>

# Course Variants

- UiO – INF3151 – Thomas Plagemann
  - <https://www.uio.no/studier/emner/matnat/ifi/INF3151/v19/timeplan/index.html#FOR>
- Princeton – COS 318 – Jaswinder P. Singh
  - <http://www.cs.princeton.edu/courses/archive/fall18/cos318/schedule.html>