

Processes and Non-Preemptive Scheduling

Tore Larsen
University of Tromsø

Otto J. Anshus
University of Tromsø, University of Oslo

Kai Li
Princeton University

An aside on concurrency

- Timing and sequencing of events are key concurrency issues
- We will study classical OS concurrency issues, including implementation and use of classic OS mechanisms to support concurrency
- In later courses on concurrency and on parallel programming you will revisit this material
- In a later course on distributed systems you may want to use formal tools to better understand and model timing and sequencing
- Emerging CPUs are typically multi core, and multi threaded
 - *Challenges and opportunities to design cost-effective high-performance systems.*

Process (I)

- An instance of a program under execution
 - Program specifying (logical) control-flow (thread)
 - Data
 - Private address space
 - Open files
 - Running environment
- The most important operating system concept
- Used for supporting the concurrent execution of independent or cooperating program instances
- Used to structure applications and systems

Processes (II)

- Classically, processes were, using today's terminology; “Single Threaded”
- Sequential program: Single process
- Parallel program: Multiple cooperating processes
 - This works fine in principle, but interaction among processes incurs significant OS time overhead.

Processes (III)

- “Modern” process: “Process” and “Thread” are separate concepts
- Process—Unit of Resource Allocation—Defines the context
- Thread—Control Thread—Unit of execution scheduling
- Every process must contain one or more threads
- Every (?) thread exists within the context of a process

User- and Kernel-Level Thread Support

- User-level threads within a process are
 - Indiscernible by OS
 - Scheduled by (user-level) scheduler in process
- Kernel-level threads
 - Maintained by OS
 - Scheduled by OS

Supporting and Using Processes

- Multiprogramming
 - Supporting concurrent execution (*parallel or transparently interleaved*) of multiple processes (or threads), often assumed to be related to different problems
 - Achieved by process- or context switching, switching the CPU(s) back and forth among the processes, keeping track of each process' progress
- Concurrent programs
 - Programs (or threads) that exploit multiprogramming for some purpose (e.g. performance, structure)
 - Operating systems is important application area for concurrent programming. Many others (event driven programs, servers, ++)
 - <http://csapp.cs.cmu.edu/public/ch12-preview.pdf>

Implementing processes

- OS needs to keep track of all processes
 - Keep track of it's progress
 - Parent process
 - Metadata (priorities etc.) used by OS
 - Memory management
 - File management
- Process table with one entry (Process Control Block) per process
- Will also align processes in *queues*

Primitives of Processes

- Creation and termination
 - `fork`, `exec`, `wait`, `kill`
- Signals
 - Action, Return, Handler
- Operations
 - `block`, `yield`
- Synchronization
 - We will talk about this later

fork (UNIX)

- Spawns a new process (with new PID)
- Called in parent process
- Returns in parent *and* child process
- Return value in parent is child's PID
- Return value in child is '0'
- Child gets duplicate, but separate, copy of parent's user-level virtual address space
- Child gets identical copy of parent's open file descriptors

fork, exec, wait, kill

- Return value tested for error, zero, or positive
- Zero, this is the child process
 - Typically redirect standard files, and
 - Call Exec to load a new program instead of the old
- Positive, this is the parent process
- Wait, parent waits for child's termination
 - Wait issued before corresponding exit: Parent blocks until exit
 - Exit issued before corresponding wait: Child becomes zombie (un-dead) until wait
- Kill, specified process terminates

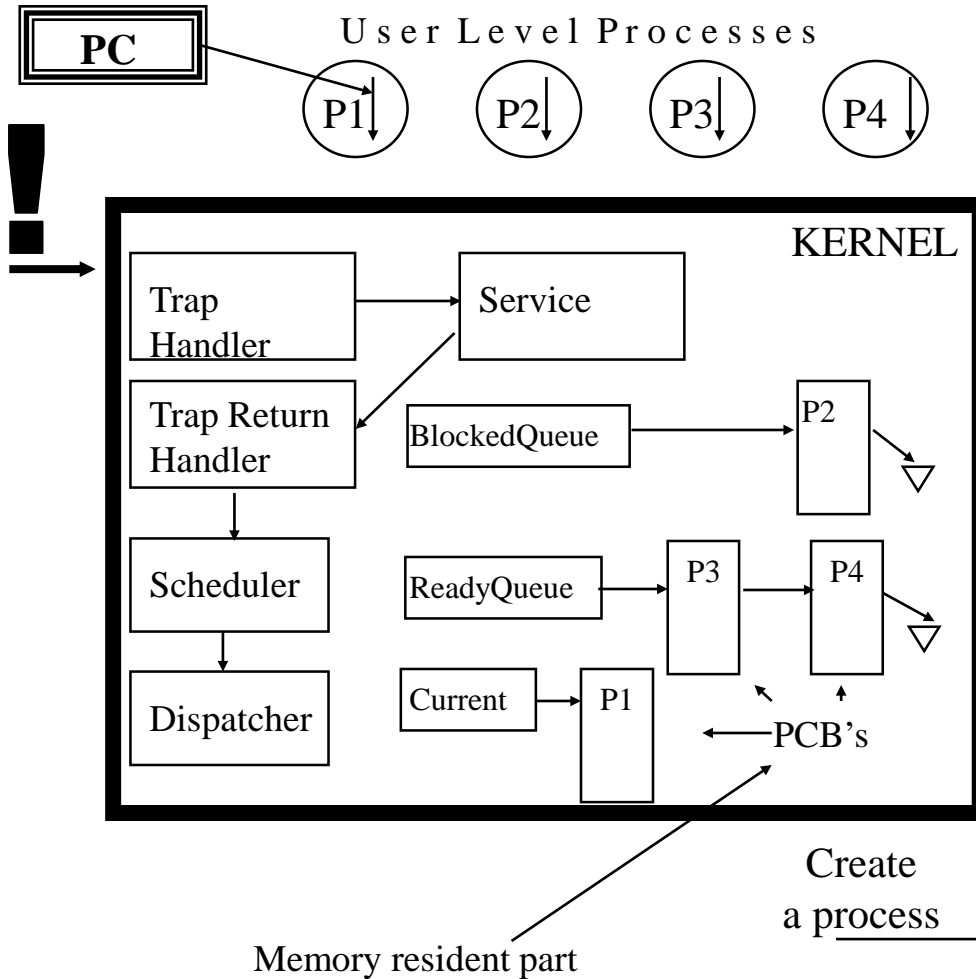
When may OS switch contexts?

- Only when OS runs
 - Events potentially causing a context switch:
 - Process created (`fork`)
 - Process exits (`exit`)
 - Process blocks implicitly (I/O, IPC)
 - Process blocks explicitly (`yield`)
 - User or System Level Trap
 - HW
 - SW: User level System Call
 - Exception
 - Kernel preempts current process
 - Potential scheduling decision at “any of above”
 - + “*Timer*” to be able to limit running time of processes
-
- The diagram consists of two labels on the right side: 'Non-Preemptive scheduling' and 'Preemptive scheduling'. From 'Non-Preemptive scheduling', an arrow points to the first four items of the list (Process created, Process exits, Process blocks implicitly, Process blocks explicitly). From 'Preemptive scheduling', an arrow points to the last two items of the list (User or System Level Trap, Kernel preempts current process).

Context Switching Issues

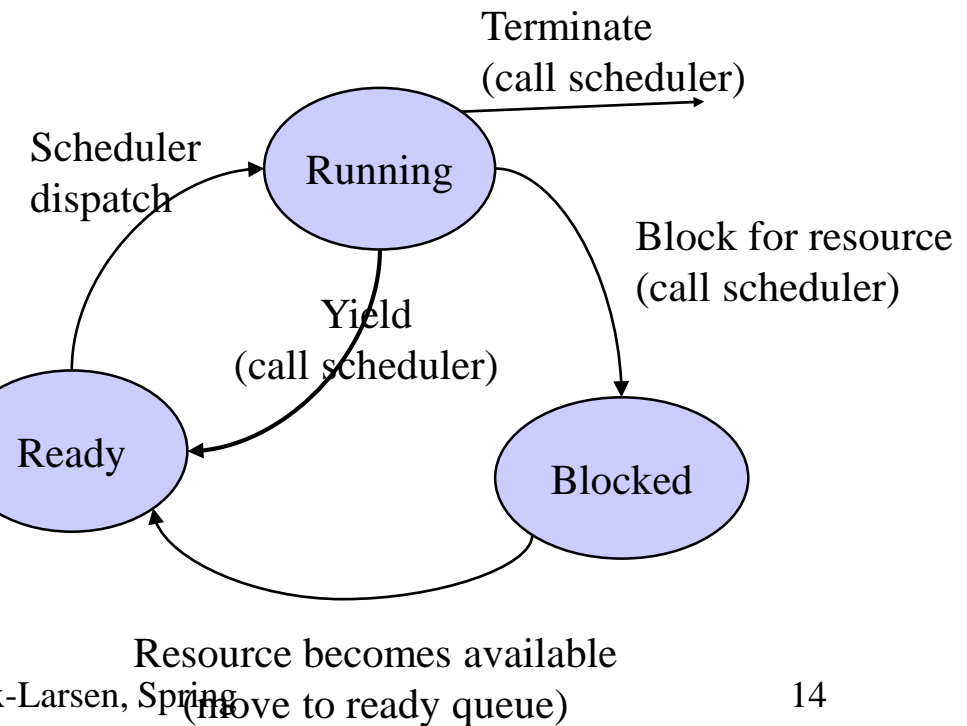
- Performance
 - Should be no more than a few microseconds
 - Most time is spent SAVING and RESTORING the context of processes
 - Less processor state to save, the better
 - Pentium has a multitasking mechanism, but SW can be faster if it saves less of the state
 - How to save time on the copying of context state?
 - Re-map (address) instead of copy (data)
- Where to store Kernel data structures “shared” by all processes
 - Memory
- How to give processes a fair share of CPU time
 - Preemptive scheduling, time-slice defines maximum time interval between scheduling decisions

Example Process State Transitions

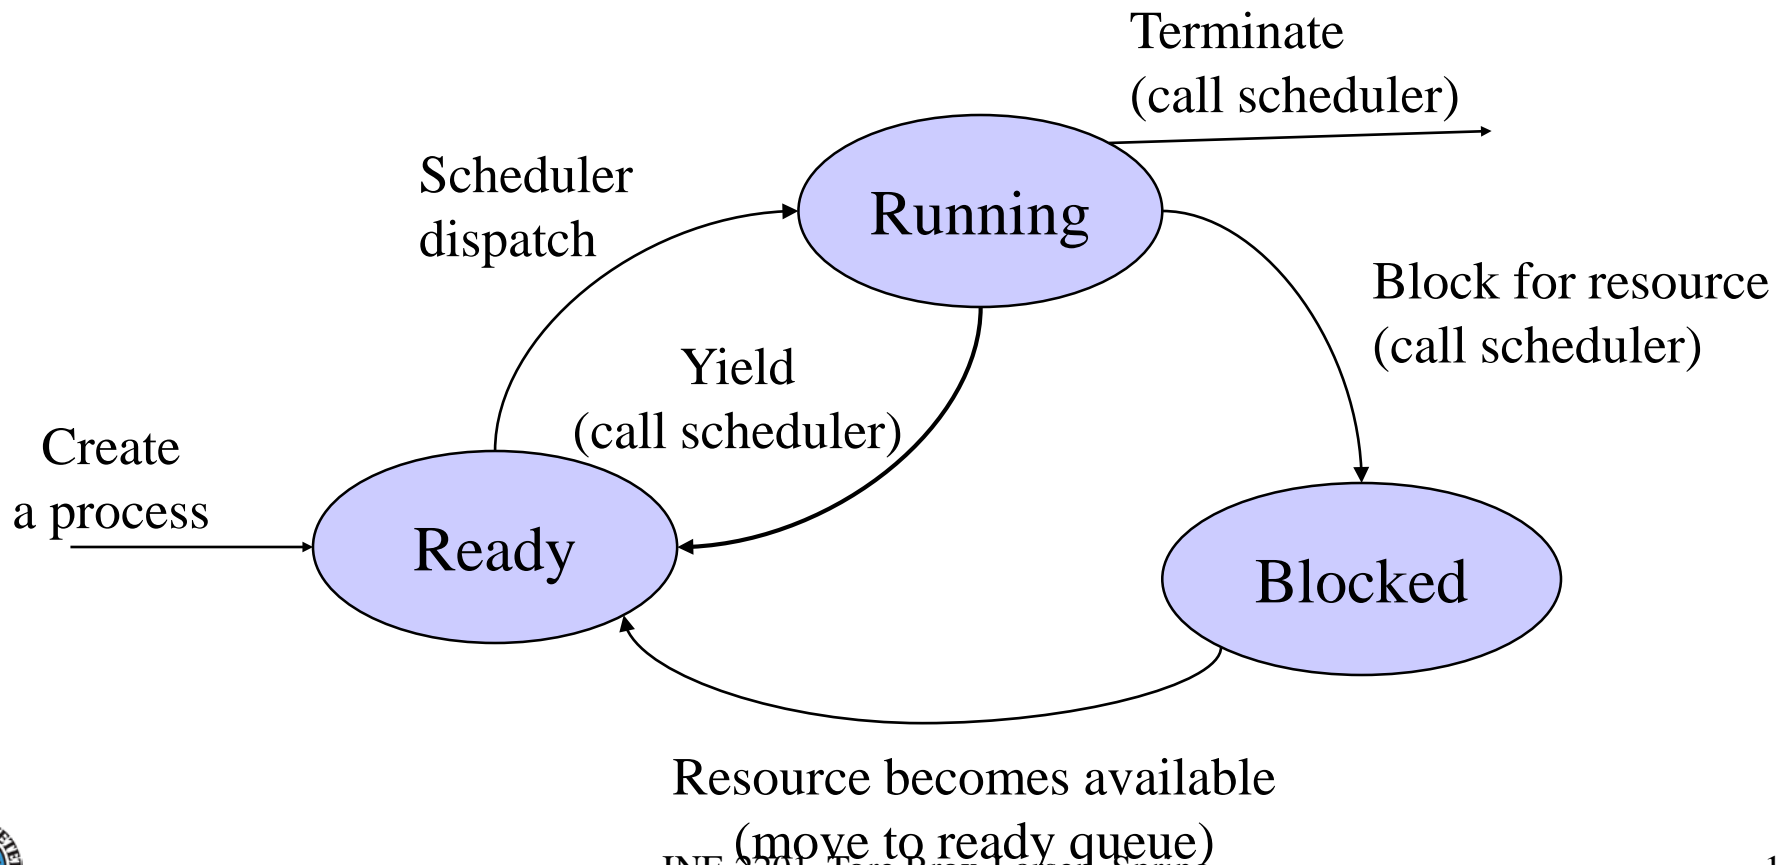


MULTIPROGRAMMING

- Uniprocessor: *Interleaving* ("pseudoparallelism")
- Multiprocessor: *Overlapping* ("true parallelism")



Process State Transition of Non-Preemptive Scheduling



Scheduler

- Non-preemptive scheduler invoked by explicit block or yield calls, possibly also fork and exit
- The simplest form

Scheduler:

save current process state (store to PCB)

choose next process to run

dispatch (load PCB and run)

- Does this work?
 - PCB (something) must be resident in memory
 - Remember the stacks

Stacks

- Remember: *We have only one copy of the Kernel in memory*
 - => all processes “execute” the same kernel code
 - => Must have a kernel stack for each process
- Used for storing parameters, return address, locally created variables in *frames* or *activation records*
- Each process
 - user stack
 - kernel stack
 - always empty when process is in user mode executing instructions
- Does the Kernel need its own stack(s)?

More on Scheduler


- Should the scheduler use a special stack?
 - Yes,
 - because a user process can overflow and it would require another stack to deal with stack overflow
 - because it makes it simpler to pop and push to rebuild a process's context
 - Must have a stack when booting...
- Should the scheduler simply be a “kernel process”?
 - You can view it that way because it has a stack, code and its data structure
 - This process always runs when there is no user process
 - “Idle” process
 - In kernel or at user level?

Win NT Idle

- No runnable thread exists on the processor
 - Dispatch Idle Process (really a *thread*)
- Idle is really a dispatcher *in the kernel*
 - Enable interrupt; Receive pending interrupts; Disable interrupts;
 - Analyze interrupts; Dispatch a thread if so needed;
 - Check for deferred work; Dispatch thread if so needed;
 - Perform power management;

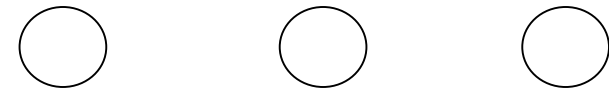
Threads and Processes

Trad. Threads

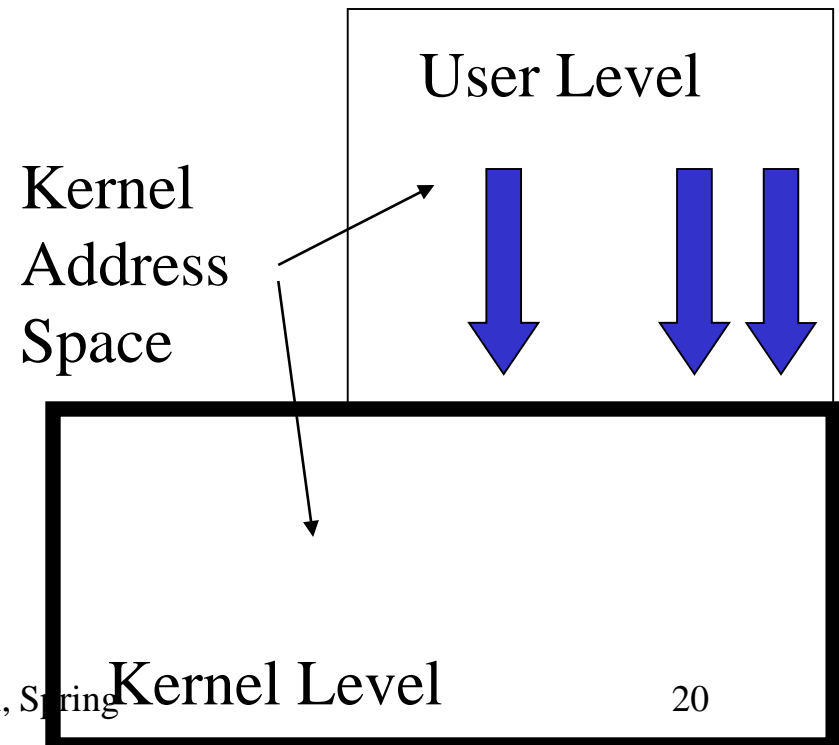


Project OpSys

Processes in individual address spaces



Kernel threads



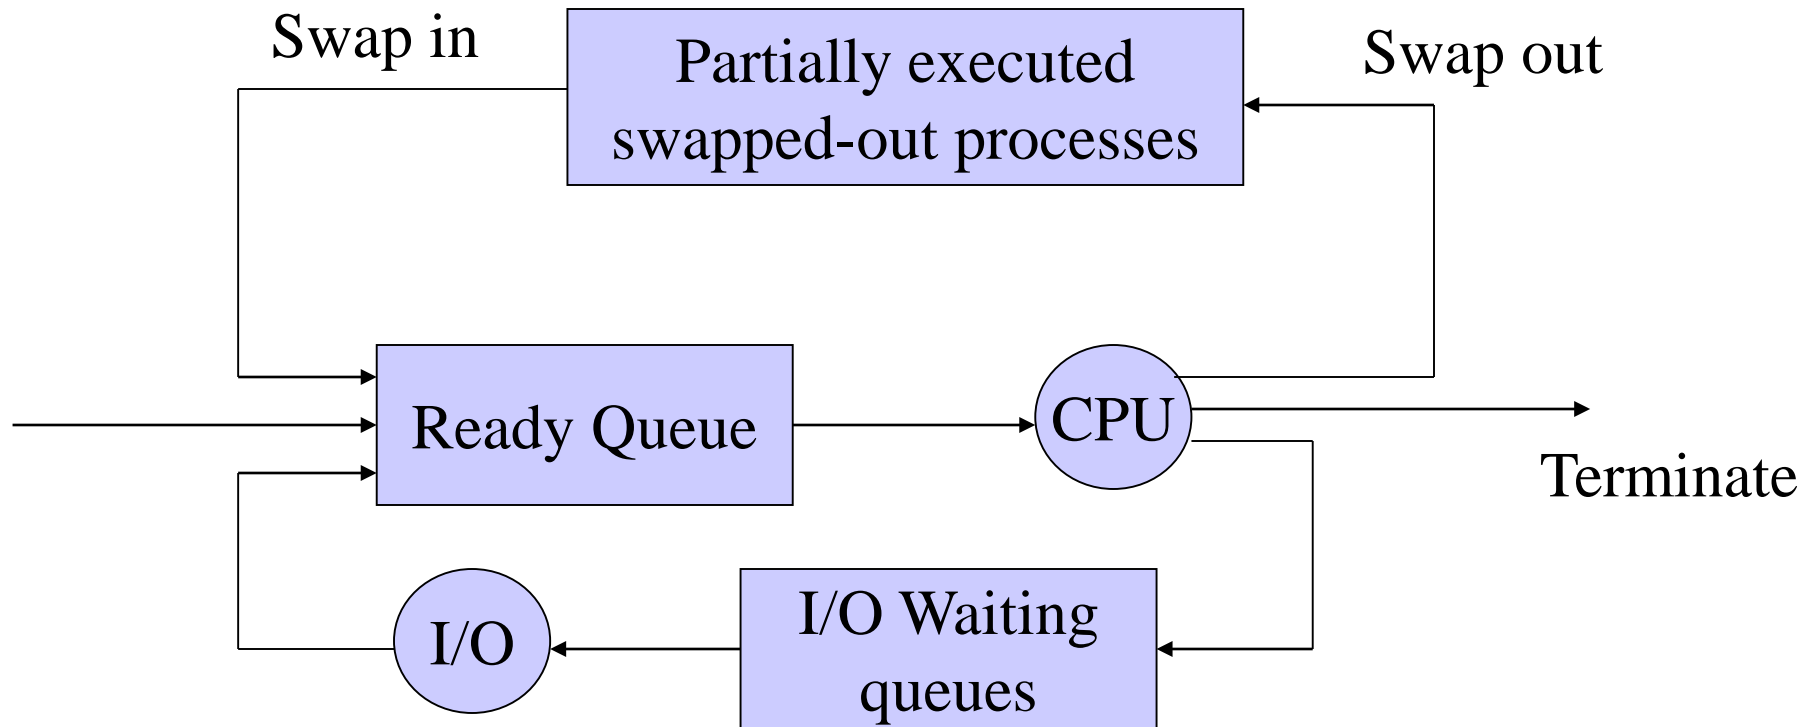
Where Should PCB Be Saved?

- Save the PCB on user stack
 - Many processors have a special instruction to do it efficiently
 - But, need to deal with the overflow problem
 - When the process terminates, the PCB vanishes
- Save the PCB on the kernel heap data structure
 - May not be as efficient as saving it on stack
 - But, it is very flexible and no other problems

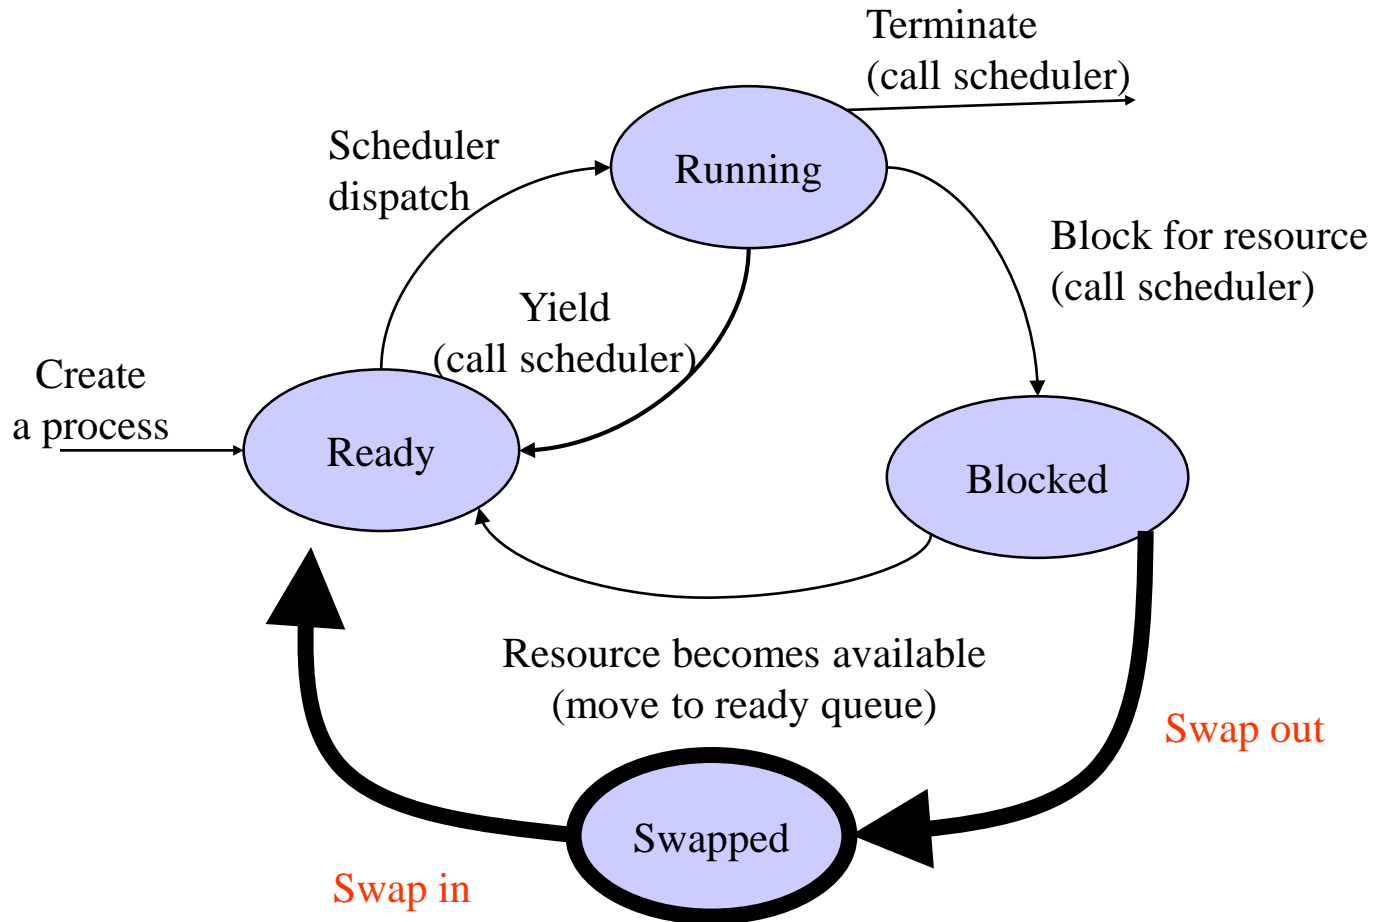
Job swapping

- The processes competing for resources may have combined demands that results in poor system performance
- Reducing the degree of multiprogramming by moving some processes to disk, and temporarily not consider them for execution may be a strategy to enhance overall system performance
 - From which states(s), to which state(s)? *Try extending the following examples using two suspended states.*
- Check&relate: MOS 3e.: Ch. 2.1.7, 3.2.1 and 3.3

Job Swapping, ii



Add Job Swapping to State Transition Diagram



Concurrent Programming w/ Processes

- Clean programming model
 - File tables are shared
 - User address space is private
 - Processes are protected from each other
 - Sharing requires some sort of IPC (InterProcess Communication)
- Slower execution
 - Process switch, process control expensive
 - IPC expensive

I/O Multiplexing:

More than one State Machine per Process

- `select` blocks for any of multiple specified I/O events
- Handle (one of the events) that unblocks `select`
 - Advance appropriate state machine
- Repeat
- See <http://csapp.cs.cmu.edu/public/ch12-preview.pdf>

Concurrent prog. w/ I/O Multiplexing

- Establishes several control flows (state machines) in one process
- Uses `select`
- Offers application programmer more control than process model
- Easy sharing of data among state machines
- More efficient (no process switch to switch between control flows in same process)
- Possibly trickier programming