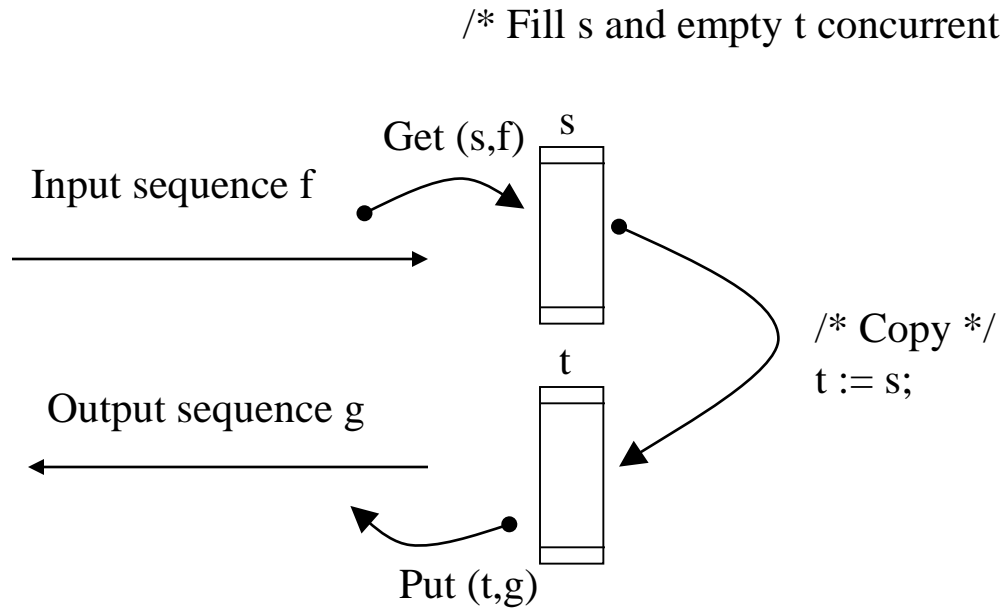


# Semaphores

Tore Larsen, UiT,  
Otto J. Anshus, UiT, UiO

# Concurrency: Double buffering



Get(s,f);

Repeat

**Copy;**

cobegin

**Put(t,g);**

**Get(s,f);**


coend;

until completed;

Specifies  
concurrent  
execution

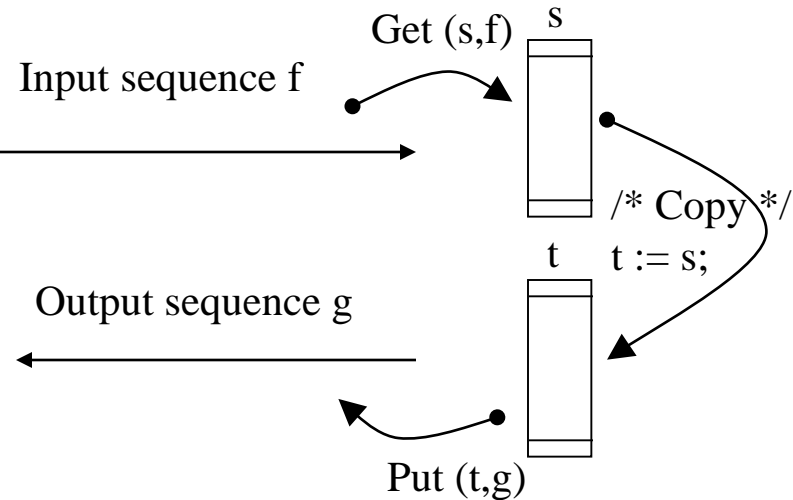
(Threads)

- **Put and Get** are disjunct
- ... but not with regards to **Copy**!



# Concurrency: Double buffering

/\* Fill s and empty t **concurrently**: OS Kernel will do preemptive scheduling of GET, COPY and PUT\*/



**Three threads executing concurrently:**

{put\_thread||get\_thread||copy\_thread} /\* Assume preemptive scheduling by kernel \*/

**Proposed code:**



**copy\_thread::** \*{acq(lock\_t); acq(lock\_s); **t=f**; rel(lock\_s); rel(lock\_t);}

**get\_thread::** \*{ack(lock\_s); **s=f**; rel(lock\_s);}

**put\_thread::** \*{ack(lock\_t): **g=t**; rel(lock\_t);}

•**Not bad, but NO ORDER**

•And as Thomas once said at the beginning of the course a few years back:  
**Ordnung Muss Sein!**



# Protecting a Shared Variable

- Remember: we need a shared address space
  - threads inside a process share adr. space
- ***Acquire(mutex); count++; Release(mutex);***
- **(1) Acquire(mutex) system call**
  - User level library
    - **(2) Push parameters onto stack**
    - **(3) Trap to kernel (int instruction)**
  - Kernel level
    - Int handler
      - **(4) Verify valid pointer to mutex**
      - Jump to code for Acquire()
        - **(5) mutex closed: block caller: insert(current, mutex\_queue)**
        - **(6) mutex open: get lock**
    - User level: **(7) execute count++**
  - **(8) Release(mutex) system call**

# Issues

- How “long” is the critical section?
- Competition for a mutex/lock
  - Uncontended = rarely in use by someone else
  - Contended = often used by someone else
  - Held = currently in use by someone
- Think about the results of these options
  - Spinning on low-cont. lock
  - Spinning on high-cont. lock
  - Blocking on low-cont. lock
  - Blocking on high-cont. lock

# Block/unblock syscalls

- Block
  - Sleep on token
- Unblock
  - Wakes up first sleeper
- By the way
  - Remember that “test and set” works both at user and kernel level

# Implementing Block and Unblock

- Block(lock)
  - Spin on lock.guard
  - Save context to TCB
  - Enqueue TCB
  - Clear spin lock.guard
  - goto scheduler
- Unblock(lock)
  - Spin on lock.guard
  - Dequeue a TCB
  - Put TCB in ready\_queue
  - Clear spin lock.guard


# Two Kinds of Synchronization

Threads inside one process: Shared address space. They can access the same variables

Process w/two threads


## MUTEX

Acquire will let first caller through, and then block next until Release



## CONDITION SYNCHRONIZATION

Acquire will block first caller until Release





# Think about ...

- Mutual exclusion using Acquire - Release:
  - Easy to forget one of them
  - Difficult to debug. must check all threads for correct use: “Acquire-CR-Release”
  - No help from the compiler?
    - It does not understand that we mean to say MUTEX
    - But could
      - check to see if we always match them “left-right”
      - associating a variable with a Mutex, and never allow access to the variable outside of CR

# Semaphores (Dijkstra, 1965)

Published as an appendix to the paper on the T.H.E. operating system

- “Down(s)”/“Wait(s)”/“P(s)”
  - Atomic
  - DELAY (block, or busy wait) if not positive
  - Decrement semaphore value by 1
- “Up(s)”/”Signal(s)”/ “V(s)”
  - Atomic
  - Increment semaphore by 1
  - Wake up a waiting thread *if any*

```
P(s) {  
    if (--s < 0)  
        Block(s);  
}
```

MUTEX




```
V(s) {  
    if (++s <= 0)  
        Unblock(s);  
}
```

Can get negative s: counts number of waiting threads

s is NOT accessible through other means than calling P and V

# An aside on Dijkstra

- Dutch, moved to UT/austin
- 1972 Turing Award Winner
- [Go to statement considered harmful](#)
- [Homepage](#)
- EDSAC [Summer School](#)



## Semaphores can be used for ...

- Mutual exclusion (solution of critical section problem).  
Binary semaphore
- Resources with multiple instances (e.g. buffer slots in producer/consumer problem. Counting semaphore
- Signaling events

# Examples of classic synchronization problems

- Critical Section
- Producer/Consumer
- Reader/Writer
- Sleeping Barber
- Dining Philosophers

# Semaphores w/Busy Wait

P(s):

```
while (s <= 0) { };  
s--;
```

V(s):

```
s++;
```

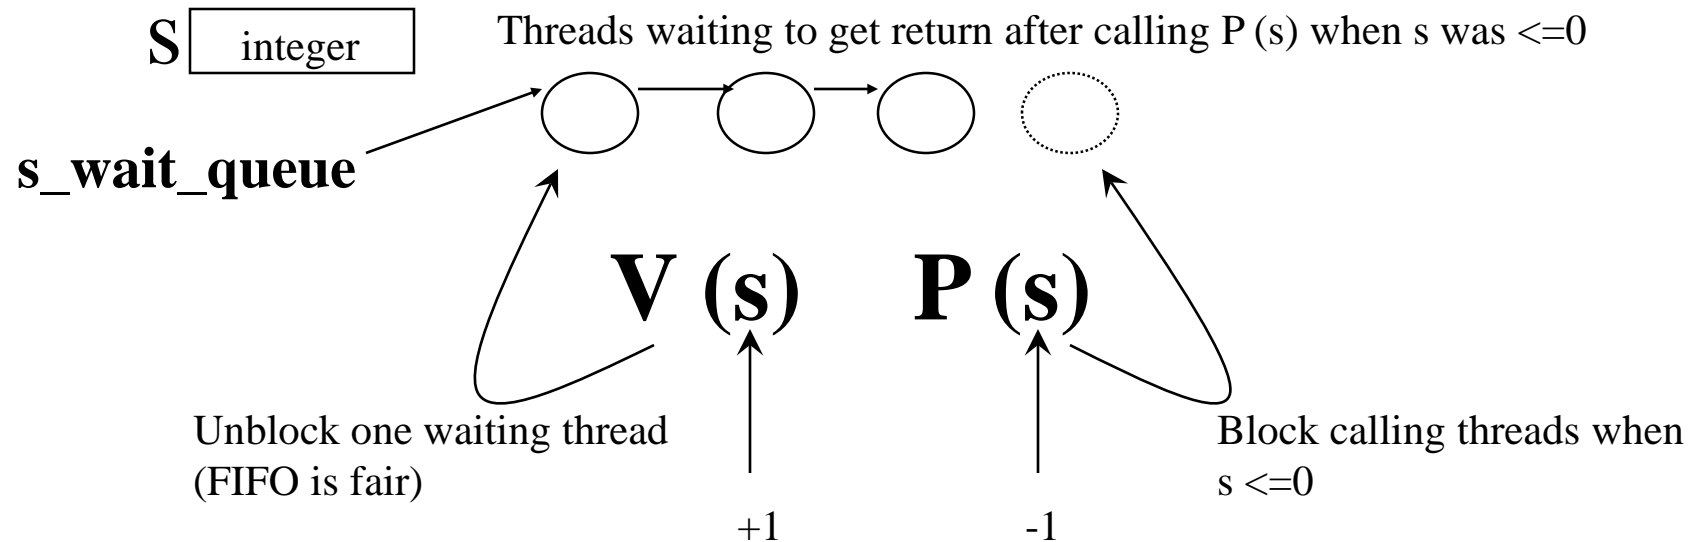
ATOMIC  
(NB: mutex around  
while can create a  
problem...)

*If spinning inside mutex  
V will not get in:*  
• *Must open mutex, say,  
between every iteration  
of while to make it possible  
for V to get in*

- Starvation possible (in theory)?
- Does it matter in practise?

- *Costly*
- *Starvation possible*
  - *Of P's*
  - *Of V's*


# The Structure of a Blocking Semaphore Implementation



- Atomic: Disable interrupts
- Atomic: P() and V() as System calls
- Atomic: Entry-Exit protocols


# Using Semaphores

**“The Signal”**




NB: remember to set the initial semaphore value!

**“The Mutex”**




**“The Team”**





# Simple to debug?



THEY ARE FOREVER WAITING FOR EACH OTHERS SIGNAL

# Semaphores w/Busy Wait

P: Passieren == to pass

P: Proberen == to test

Dutch words

V: Vrijmagen == to make free

V: Verhogen == to increment

P(s):

```
while (s <= 0) { };  
s--;
```

V(s):

```
s++;
```

mutex




P == wait == down, V == signal == up


Why so many names?

- Down, up: what the ops *do*
- Wait, signal: what the ops are *used for*
- P, V: the original names by Dijkstra

# Rendezvous between two threads



# Bounded Buffer using Semaphores



**PUT (msg):**

```
P(nonfull);
P(mutex);
<insert>
V(mutex);
V(nonempty);
```

**GET (buf):**

```
P(nonempty);
P(mutex);
<remove>
V(mutex);
V(nonfull);
```

**Rules for the buffer B:**

•No Get when empty

•No Put when full

•B shared, so must have  
mutex between Put and  
Get

**Use one semaphore for  
*each condition* we must  
wait for to become TRUE:**

•B empty: nonempty:=0;


•B full: nonfull:=N

•B mutex: mutex:=1;

•Is Mutex needed when only 1 P and 1 C?

•PUT at one end, GET at other end

# Semaphores w/Busy Wait



**If P spinning inside mutex then V will not get in**

- Must *open mutex*, say, between every iteration of while to make it possible for V to get in
  - Costly
    - Every 10th iteration?
    - latency
- Starvation possible, Lady Luck may ignore some threads
  - Of P's
  - Of V's

# Hard life...

- Implementing the P and V of semaphores
  - If WAIT is done by blocking
    - Expensive
    - Must open mutex
      - But no logical issues since we now have a waiting queue and will not get starvation
  - If done by spinning
    - Must open mutex during spin to let V in
      - Starvation of P's and V's possible
        - May not be a problem in practise
- What can a poor (perhaps somewhat theoretical oriented) Computer Scientist do?
  - Research (“I can do better”)
  - Publish (So other people can say “I can do better”)



# Implementing Semaphores w/mutex

```
P(s) {  
    Acquire(s.mutex);  
    if (--s.value < 0) {  
        Release(s.mutex);  
        Acquire(s.delay);  
    } else  
        Release(s.mutex);  
}  
  
V(s) {  
    Acquire(s.mutex);  
    if (++s.value <= 0)  
        Release(s.delay);  
    Release(s.mutex);  
}
```

## ◆ Kotulski (1988)

- Two processes call P(s) (s.value is initialized to 0) and preempted after Release(s.mutex)
- Two other processes call V(s)

# Hemmendinger's solution (1988)

```
P(s) {
    Acquire(s.mutex);
    if (--s.value < 0) {
        Release(s.mutex);
        Acquire(s.delay);
    }
    Release(s.mutex);
}

V(s) {
    Acquire(s.mutex);
    if (++s.value <= 0)
        Release(s.delay);
    else
        Release(s.mutex);
}
```

- ◆ The idea is not to release s.mutex and turn it over individually to the waiting process
- ◆ P and V are executing in locksteps



# Kearn's Solution (1988)

```
P(s) {
    Acquire(s.mutex);
    if (--s.value < 0) {
        Release(s.mutex);
        Acquire(s.delay);
        Acquire(s.mutex);
        if (--s.wakecount > 0)
            Release(s.delay);
    }
    Release(s.mutex);
}

V(s) {
    Acquire(s.mutex);
    if (++s.value <= 0) {
        s.wakecount++;
        Release(s.delay);
    }
    Release(s.mutex);
}
```

Two Release( s.delay) calls are also possible

# Hemmendinger's Correction (1989)

```
P(s) {  
    Acquire(s.mutex);  
    if (--s.value < 0) {  
        Release(s.mutex);  
        Acquire(s.delay);  
        Acquire(s.mutex);  
        if (--s.wakecount > 0)  
            Release(s.delay);  
    }  
    Release(s.mutex);  
}
```

```
V(s) {  
    Acquire(s.mutex);  
    if (++s.value <= 0) {  
        s.wakecount++;  
        if (s.wakecount == 1)  
            Release(s.delay);  
    }  
    Release(s.mutex);  
}
```

Correct but a complex solution


# Hsieh's Solution (1989)

```
P(s) {  
    Acquire(s.delay);  
    Acquire(s.mutex);  
    if (--s.value > 0)  
        Release(s.delay);  
    Release(s.mutex);  
}
```

```
V(s) {  
    Acquire(s.mutex);  
    if (++s.value == 1)  
        Release(s.delay);  
    Release(s.mutex);  
}
```

- ◆ Use Acquire(s.delay) to block processes
- ◆ Correct but still a constrained implementation


# Dining Philosophers



- Each: need 2 forks to eat
- 5 philosophers: 10 forks
- 5 forks: 2 can eat concurrently

## Things to observe:

- A fork can only be used by one at a time
- No deadlock, please
- No starving, please
- Concurrent eating, please


**s**  
  
 s(i): One semaphore per fork to be used in **mutex** style P-V

**Mutex on whole table:**  $P(\text{mutex});$   $T_i$   
 • *I can eat at a time* eat;  
 $V(\text{mutex});$

**Get L; Get R;**  $P(s(i));$   $T_i$   
 • *Deadlock possible*  $P(s(i+1));$   
 eat;  
 $S(i) = 1$  initially  $V(s(i+1));$   
 $V(s(i));$

**Get L; Get R if free else Put L;**  $T_i$   
 • *Starvation possible*


# Dining Philosophers



To avoid starvation they could look after each other:

- Entry:** If L and R is not eating I can
- Exit:** If L (R) wants to eat and L.L (R.R) is not eating I start him eating

One semaphore per philosopher  
Used in **signal** style



- Thinking
- Eating
- Want

$S(i) = 0$  initially

$T_i$

```
While (1) {
  <think>
  ENTRY;
  <eat>
  EXIT;
}
```

```
P(mutex);
state(i):=Want;
if (state(i-1) !=Eating AND state(i+1) != Eating)
  /*Safe to eat*/
  state(i):=Eating;
  V(s(i)); /*Because */ }
V(mutex);
P(s(i)); /*Init was 0!! I or neighbor must say V(i) to myself!*/
```

```
P(mutex);
state(i):=Thinking;
if (state(i-1)=Want AND state(i-2) !=Eating)
{
  state(i-1):=Eating;
  V(s(i-1)); /*Start Left neighbor*/
}
/*Analogue for Right neighbor*/
V(mutex);
```

Trouble: **starvation** pattern possible:

2&4 at table, 1&3 hungry

2 gets up, 1 sits down

4 gets up, 3 sits down

3 gets up, 4 sits down

1 gets up, 2 sits down

Ad infinitum => Phil 0 will starve

# Last solution has a problem

Trouble in Tanenbaums solution:

**starvation** pattern possible:

2&4 at table, 1&3 hungry


2 gets up, 1 sits down

4 gets up, 3 sits down


3 gets up, 4 sits down

1 gets up, 2 sits down

Ad infinitum => Phil 0 will starve



# Dining Philosophers




Can we in a simple way do better than this one?

**Get L; Get R;**

•Deadlock possible

```
P(s(i));
P(s(i+1));
eat;
V(s(i+1));
V(s(i));
```

•Non-symmetric solution. Still quite elegant



$S(i) = 1$  initially

$T_1, T_2, T_3, T_4:$

```
P(s(i));
P(s(i+1));
<eat>
V(s(i+1));
V(s(i));
```

$T_5$

```
P(s(1));
P(s(5));
<eat>
V(s(5));
V(s((1));
```

- Remove the danger of circular waiting (deadlock)
- $T_1$ - $T_4$ : Get L; Get R;
- $T_5$ : Get R; Get L;

## Some Links

- Wikipedia: [\*Semaphore\*](#)
- Alan B. Downey: *The Little Book of Semaphores*  
[\*Book\*](#)
- Creature Mann: Santa and his Helpers: [\*Video\*](#)
- Udacity: Mutual Execution vs Synchronization: [\*Video\*](#)
- Jouni Leppäjärvi: [\*Master's Thesis\*](#)