

Monitors Condition Variables

Otto J. Anshus

University of {Tromsø, Oslo}

With some revisions by
Kai Li and Tore Brox-Larsen

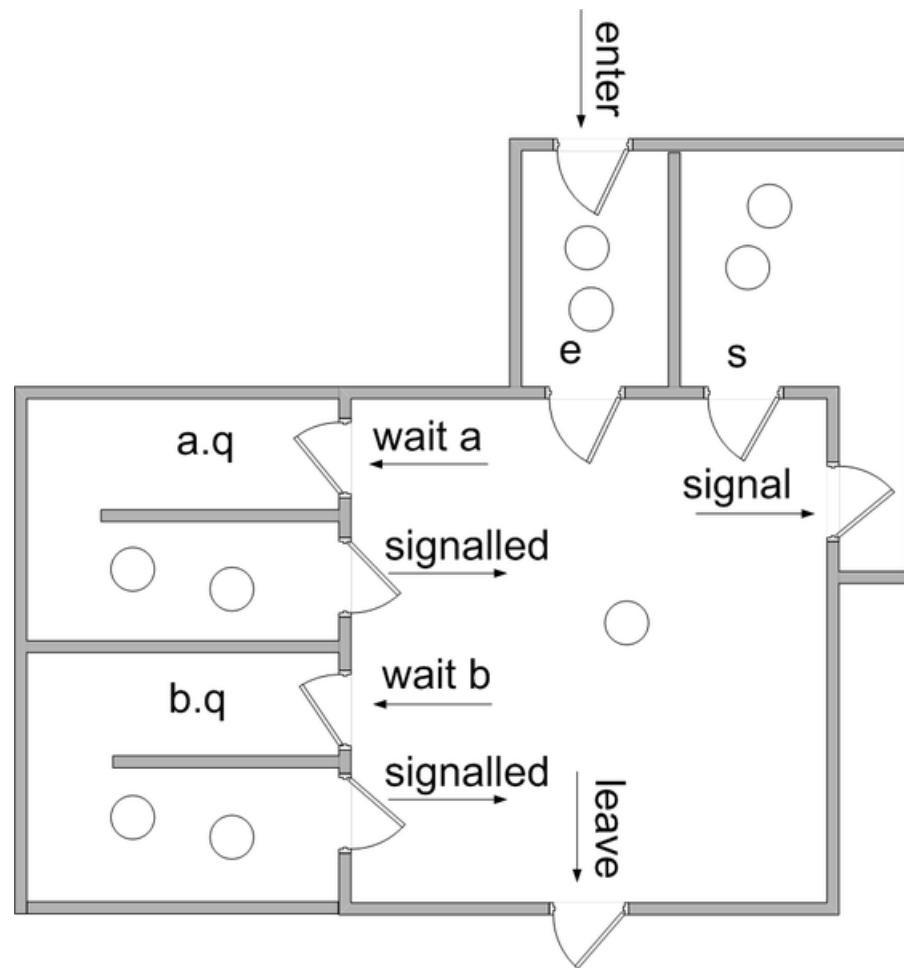
Monitor: Higher Abstraction Level Mechanism

- Idea by Brinch-Hansen 1973 in the textbook “Operating System Principles”
 - Structure an OS into a set of modules each implementing a resource scheduler
- Tony Hoare – 1974
 - Combine together in each module
 - Mutex
 - Shared data
 - Access methods to shared data
 - Condition synchronization
 - Local code and data


Basic Components

- *Monitor procedures* are **mutually exclusive**
 - called/executed by threads (monitor is implemented by all threads, data variables are shared)
 - “called” by processes (without shared address space) is also possible (HOW?)
- Condition “variables” (*declared by user*)
 - just a naming of wait queues
 - select a meaningful name (good programming practice)
- Wait (cond_var_name) (*called by monitor procedures*)
 - calling thread will unconditionally be removed as *current* and inserted into the named waiting queue
 - then the OS kernel scheduler must select another process from the ready_queue to become the new *current*
- Signal (cond_var_name) (*called by monitor procedures*)
 - resume (wakeup) a blocked thread


Blocking Condition Variables (Hoare style)



Non-Blocking (Mesa style)



Implicit Condition Variables (Java style)



The Structure of a Monitor

MUTEX
So only ONE
monitor
procedure
executes at a
time

Proposal: $P(mname); \langle mon \rangle V(name)$

The Monitor

Main Queue

Could be the mutex wait queue

Condition Queue 1

Condition Queue n

<More to come>

Signal(): { ... }

Wait(): { ... }

Local variables

Local procedure 1

Local procedure m

Shared variables

Monitor procedure 1: { ... wait(condvar); ... }

•
•

Monitor procedure k: { ... signal(condvar); ... }

Initialization executed first time the monitor starts

•After calling, threads get
blocked and are waiting to
get in and start executing
the called monitor
procedure

•Threads waiting on a condition
variable for a condition to be true
(waiting for a signal on the
condition variable)


System implementation

User implementation

•The only way to access shared
resources is by calling a monitor
procedure

•Initialization of state variables,
executed ONCE at startup of
monitor

Threads calling a
monitor procedure.
Can also be done
as “in-line” code
in each thread



Signal and Wait

- Wait (cond)
 - Insert(caller, cond_queue)
 - **Block** this instance of the monitor procedure
 - open MUTEX to get next call (for instance from the waiting queue to get in (get next call from Main_Queue))
- Signal (cond)
 - **Stop** monitor procedure calling signal (*stop??*)
 - Start first in cond_queue, or just return if empty

Implementation of the Monitor Concept

- As a primitive in a language (Mesa, Java)
- By using semaphores (in any language)
- As a thread or as a process
 - Need a way to interact with the **thread**
 - through shared variables to deliver the parameters and name of called monitor procedure
 - Need a way to interact with the **process**
 - kernel support of shared variables across address spaces
 - using another mechanism like message passing to pass parameters and name of procedure
- **At user level**
 - **use condition variables (the queues)**
 - **wait(), signal() implemented by**
 - the operating system kernel
 - a thread package (Pthreads)
 - **mutex by P-V or a special keyword like *synchronized()***

Single Resource Monitor

All threads must follow the pattern:

Reserve;

<use shared resource>

Release;

```
/*Local functions, variables*/  
<none needed>  
/*Shared variable*/  
Boolean busy;  
/*Condition variable*/  
Condition nonbusy;
```

Reserve:

```
{  
    if (busy) wait (nonbusy);  
    busy:=TRUE;  
}
```

Release:

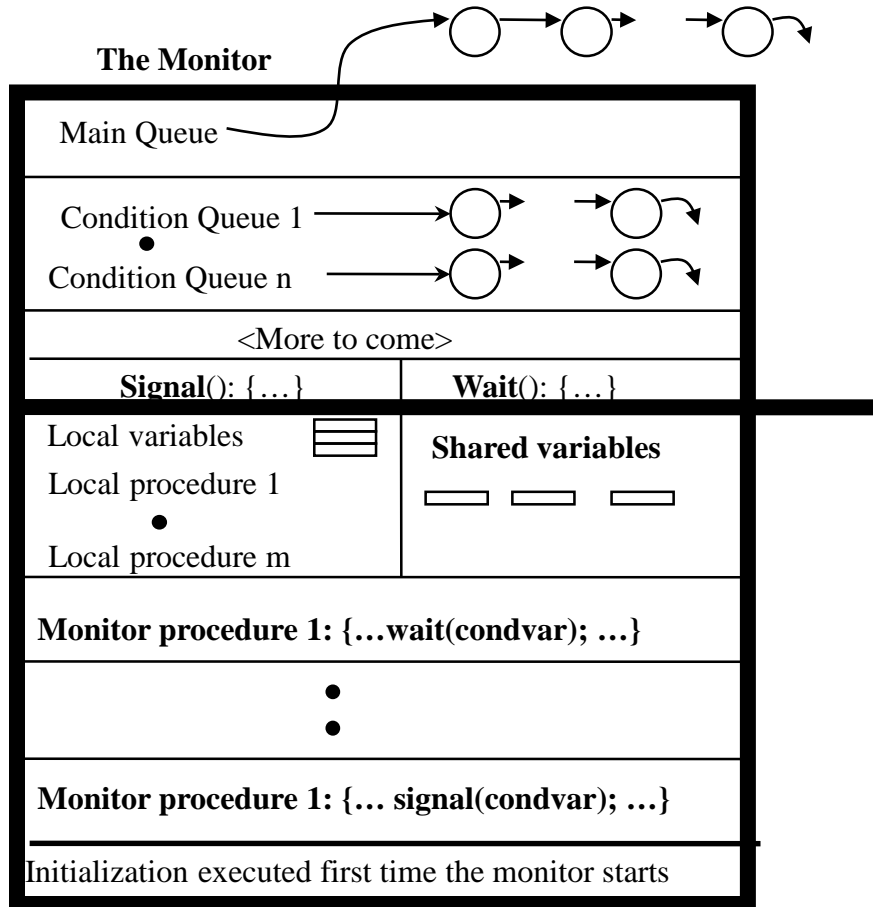
```
{  
    busy:=FALSE;  
    signal (nonbusy);  
}
```

```
/* Initialization code*/  
busy:=FALSE;  
nonbusy:=EMPTY;
```

Observe


- the shared variable
- the naming of the condition variable
- the wait and signal calls
- implements a binary semaphore ($s=0,1$)

What is a Condition Variable?



- No “value”
- Waiting queue
- Used to represent a condition we need to wait for to be TRUE
- Initial “non-value” is EMPTY :-)

Bounded Buffer Monitor



Rules for the buffer B:

- No Get when empty
 - No Put when full
 - B shared, so must have mutex between Put and Get
- One condition variable for each condition:
- nonempty
 - nonfull
 - MUTEX is already provided by the monitor

```
/*Local functions, variables*/
int in, out;
/*Shared variable*/
int B(0..n-1), count;
/*Condition variable*/
Condition nonfull, nonempty;
```

Put (int m):

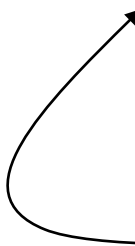
```
{ if (count=n) wait (nonfull);
  B(in):=m;
  in:=in+1 MOD n; /* MOD is % */
  count++;
  signal (nonempty); }
```

int Get:

```
{ if (count=0) wait (nonempty);
  Get:=B(out);
  out:=out+1 MOD n;
  count--;
  signal (nonfull); }
```

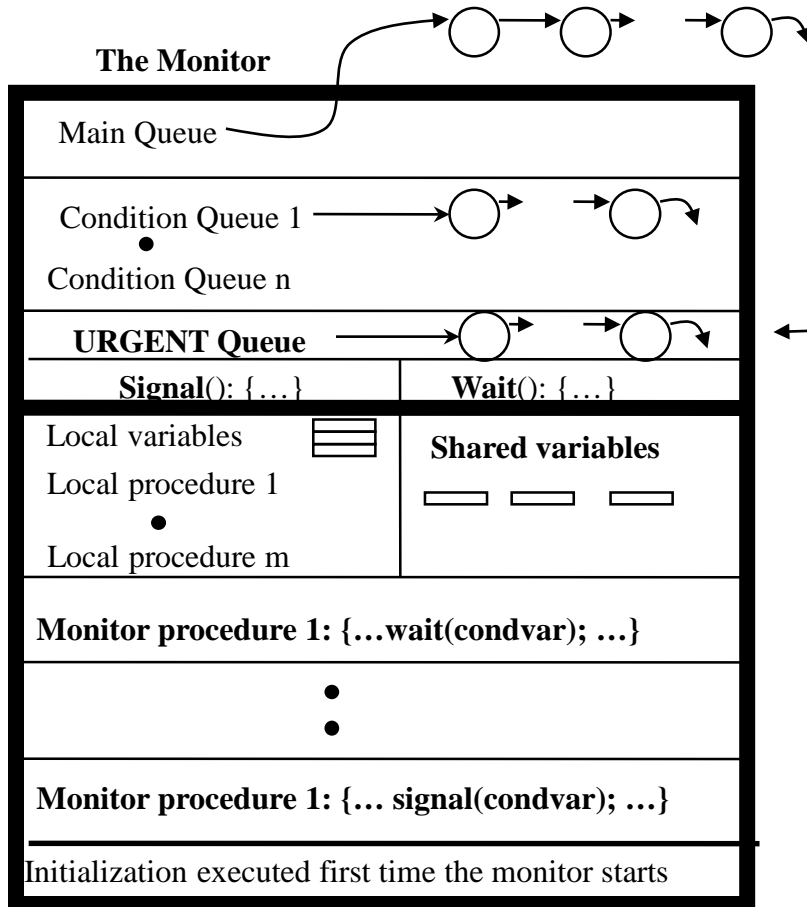
```
/* Initialization code*/
in:=out:=count:=0;
nonfull, nonempty:=EMPTY;
```

What will happen when a signal() is executed?

- Assume we have threads in Main_Queue and in a condition queue
 - Main_Queue has lower “priority” than the signaled condition queue:
 - signal() => Take first from condition queue and start it from its next instruction after the wait() which blocked it
 - The signaled thread now executes
 - ... until a wait(): block it, and take new from Main_Queue
 - ... until a signal():
 - ... until finished: take new from Main_Queue
- 

Where to allow a call to signal()?

- Look at the two monitors we have analyzed! Where is the signal() operation?
- What if we called signal somewhere else?



- The calling function instance must be blocked, awaiting return from signal()

– Need a queue for the temporary halted thread

- URGENT QUEUE

- In Hoare's monitors the signal operation must **IMMEDIATELY** start the signaled thread in order for the condition that it signals about **still to be guaranteed true** when the thread starts

Options of the Signaler

- Run the signaled monitor procedure (or thread) *immediately* (must suspend the current one right away) (**Hoare**)
 - If the signaler has other work to do, life gets complex
 - It is difficult to make sure there is nothing more to do because the signal implementation is not aware how it is used (where it is called)
 - It is easier to prove things
- Exit the monitor (**Hansen**)
 - Just let signal be the last statement before return from a monitor procedure
- Continues its execution (**Mesa**)
 - Easy to implement
 - But, the condition may not be true when the awaken process actually gets a chance to run

Condition Variables and Signal Semantics

Monitor Style	Max # Cond. Variables	Signal Semantics
Brinch-Hansen	More	Implicitly Enforces “Signal and Leave”
Hoare	More	“Signal and Block”
Mesa	More	“Signal and Continue”
Java	One (implicit)	“Signal and Continue”

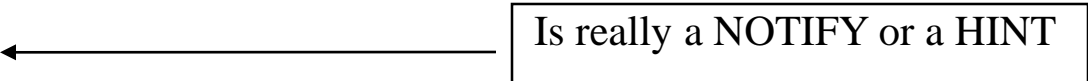
(Mutex between monitor procedures?)

- Hoare: Yes
- But not needed if we have no shared variables
 - But **signal** and **wait** must be atomic because they can access the same condition variable
 - So no gain?
 - Finer granularity (is good)
 - Makes life harder (is bad)
- Should be possible to Put and Get at each end of a buffer?
 - Try it

Performance problems of Monitors?

- Getting in through Main_Queue
 - Many can be in Main_Queue and in a condition queue waiting for a thread to execute a monitor procedure calling a signal.
 - Can take a long time before the signaler gets in
 - Need one Wait_Main_Queue and one Signal_Main_Queue?
 - But difficult when all procedures call both wait and signal
- The monitor is a potential bottleneck (“Bottleneck OS”? :))
 - Use several to avoid hot spots
- Signal must start the signaled thread immediately, so must switch thread context and save our own
 - Can have nested calls
 - Even worse for process context switches
 - Solution?
 - **Avoid starting the signaled thread immediately**
 - But then race conditions can happen so must be careful


Mesa Style “Monitor” (Birrell’s Paper)

- Condition variables are always associated with a mutex
- Wait(**mutex**, condition)
 - Atomically unlock the mutex and enqueue on the condition variable (block the thread)
 - Re-lock the lock when it is awoken
- Signal(condition) 

Is really a NOTIFY or a HINT

 - No-op if there is no thread blocked on the condition variable
 - Wake up at **some** convenient time **at least one** (if there are threads blocked)
- Broadcast(condition)
 - Wake up **all** threads waiting on the condition

Bounded Buffer Mesa Monitors



Rules for the buffer B:

- No Get when empty
- No Put when full
- B shared, so must have mutex between Put and Get

One condition for each condition:

- nonempty
- nonfull
- MUTEX is locked by LOCK and unlocked by Wait

```
/*Local functions, variables*/
int in, out, count;
/*Shared variable*/
int B(0..n-1);
/* Mutex */
mutex_t bb_mutex;
/*Condition variable*/
Condition nonfull, nonempty;
```

Put (int m):

```
LOCK bb_mutex {
    { while (count=n) wait (bb_mutex, nonfull);
      B(in):=m;
      in:=in+1 MOD n;
      count++;
      signal (nonempty); }
}
```

Wait will
UNLOCK

int Get:

```
LOCK bb_mutex {
    { while (count=0) wait (bb_mutex, nonempty);
      Get:=B(out);
      out:=out+1 MOD n;
      count--;
      signal (nonfull); }
}
```

```
/* Initialization code*/
```

```
in:=out:=count:=0;
nonfull, nonempty:=EMPTY;
```

Mesa-Style vs. Hoare-Style Monitor

- Mesa-style
 - Signaler keeps lock and CPU
 - Waiter simply put on ready queue, with no special priority
 - Must then spin and reevaluate!
 - No costly context switches immediately
 - No constraints on when the waiting thread/process must run after a “signal”
 - Simple to introduce a broadcast: wake up all
 - Good when one thread frees resources, but does not know which other thread can use them (“who can use j bytes of memory?”)
 - Can easily introduce a watch dog timer: if timeout then insert waiter in Ready_Queue and let waiter reevaluate
 - Will guard a little against bugs in other signaling processes/threads causing starvation because of a “lost” signal
- Hoare-style
 - Signaler gives up lock and waiter runs immediately
 - Waiter (now executing) gives lock and CPU back to signaler when it exits critical section or if it waits again

Instead of LOCK and UNLOCK...

```
static count = 0;  
static Cond full, empty;  
static Mutex lock;
```

```
Enter(Item item) {  
    Acquire(lock);  
    while (count==N)  
        Wait(lock, full);  
    insert item into buffer  
    count++;  
    if (count==1)  
        Signal(empty);  
    Release(lock);  
}
```

```
Remove(Item item) {  
    Acquire(lock);  
    while (!count)  
        Wait(lock, empty);  
    remove item from buffer  
    count--;  
    if (count==N-1)  
        Signal(full);  
    Release(lock);  
}
```

Can we replace “while” with “if?”

Think about the performance benefit of this solution

Programming Idiom

◆ Waiting for a resource

```
Acquire(mutex);  
while (no resource)  
    wait(mutex, cond);  
    use the resource  
Release(mutex);
```

◆ Make resource available

```
Acquire(mutex);  
    make resource  
Signal(cond);  
Release(mutex);
```

Implementing Semaphores with Mesa-Monitors

P(s)

```
{  
    Acquire( s.mutex );  
    --s.value;  
    if (s.value < 0 )  
        wait( s.mutex, s.cond );  
    Release( s.mutex);  
}
```

V(s)

```
{  
    Acquire( s.mutex );  
    ++s.value;  
    if (s.value >= 0 )  
        signal( s.cond );  
    Release( s.mutex);  
}
```

Assume that Signal wakes up exactly one awaiting thread.

Semaphore vs. Monitor

Semaphore

P(s) means WAIT if $s=0$
And $s--$

V(s) means start a waiting thread
and REMEMBER that a V call
was made: $s++$

Assume $s=0$ when V(s) is called:
If there is no thread to start this
time, the next thread to call P(s)
will get through P(s)

Monitor

Wait(cond) means unconditional WAIT

Signal(cond) means start a
waiting thread. But no memory!

Assume that the condition queue
is empty when signal() is called.
The next thread to call
Wait(cond) (by executing a
monitor procedure!) will block
because the signal() operation did
not leave any trace of the fact that
it was executed on an empty
condition waiting queue.

Equivalence

- Semaphores
 - Good for signaling
 - Not good for mutex because it is easy to introduce a bug
- Monitors
 - Good for scheduling and mutex
 - Too (maybe?) costly for simple signaling

History of Monitors

- [Brinch-Hansen](#) ('73) & [Tony Hoare](#) ('74)
 - Concept – No implementation
 - Requires Signal to be last action (Hansen)
 - Requires relinquishing CPU to signaler (Hoare)
- [Mesa](#) language ('77)
 - Monitor in language, but signaler keeps mutex & CPU
 - Waiter simply put on ready-queue
- [Modula 2+](#) ('84) & [Modula 3](#) ('88)
 - Explicit LOCK primitive
 - Mesa style monitor
- [Pthreads](#) ('95)
 - Started standardization ~'89
 - Defined by ANSI/POSIX Runtime Library
- [Java threads](#)
 - [Gosling](#), early '90s without threads
 - Use most of Pthreads primitives