

Deadlocks

INF-2201 Operating Systems Fundamentals – Spring 2015

Bård Fjukstad bard.fjukstad@uit.no b.fjukstad@met.no

Based on a presentations created by
Daniel Stødle, NORUT
And Kai Li and Andy Bavier, Princeton

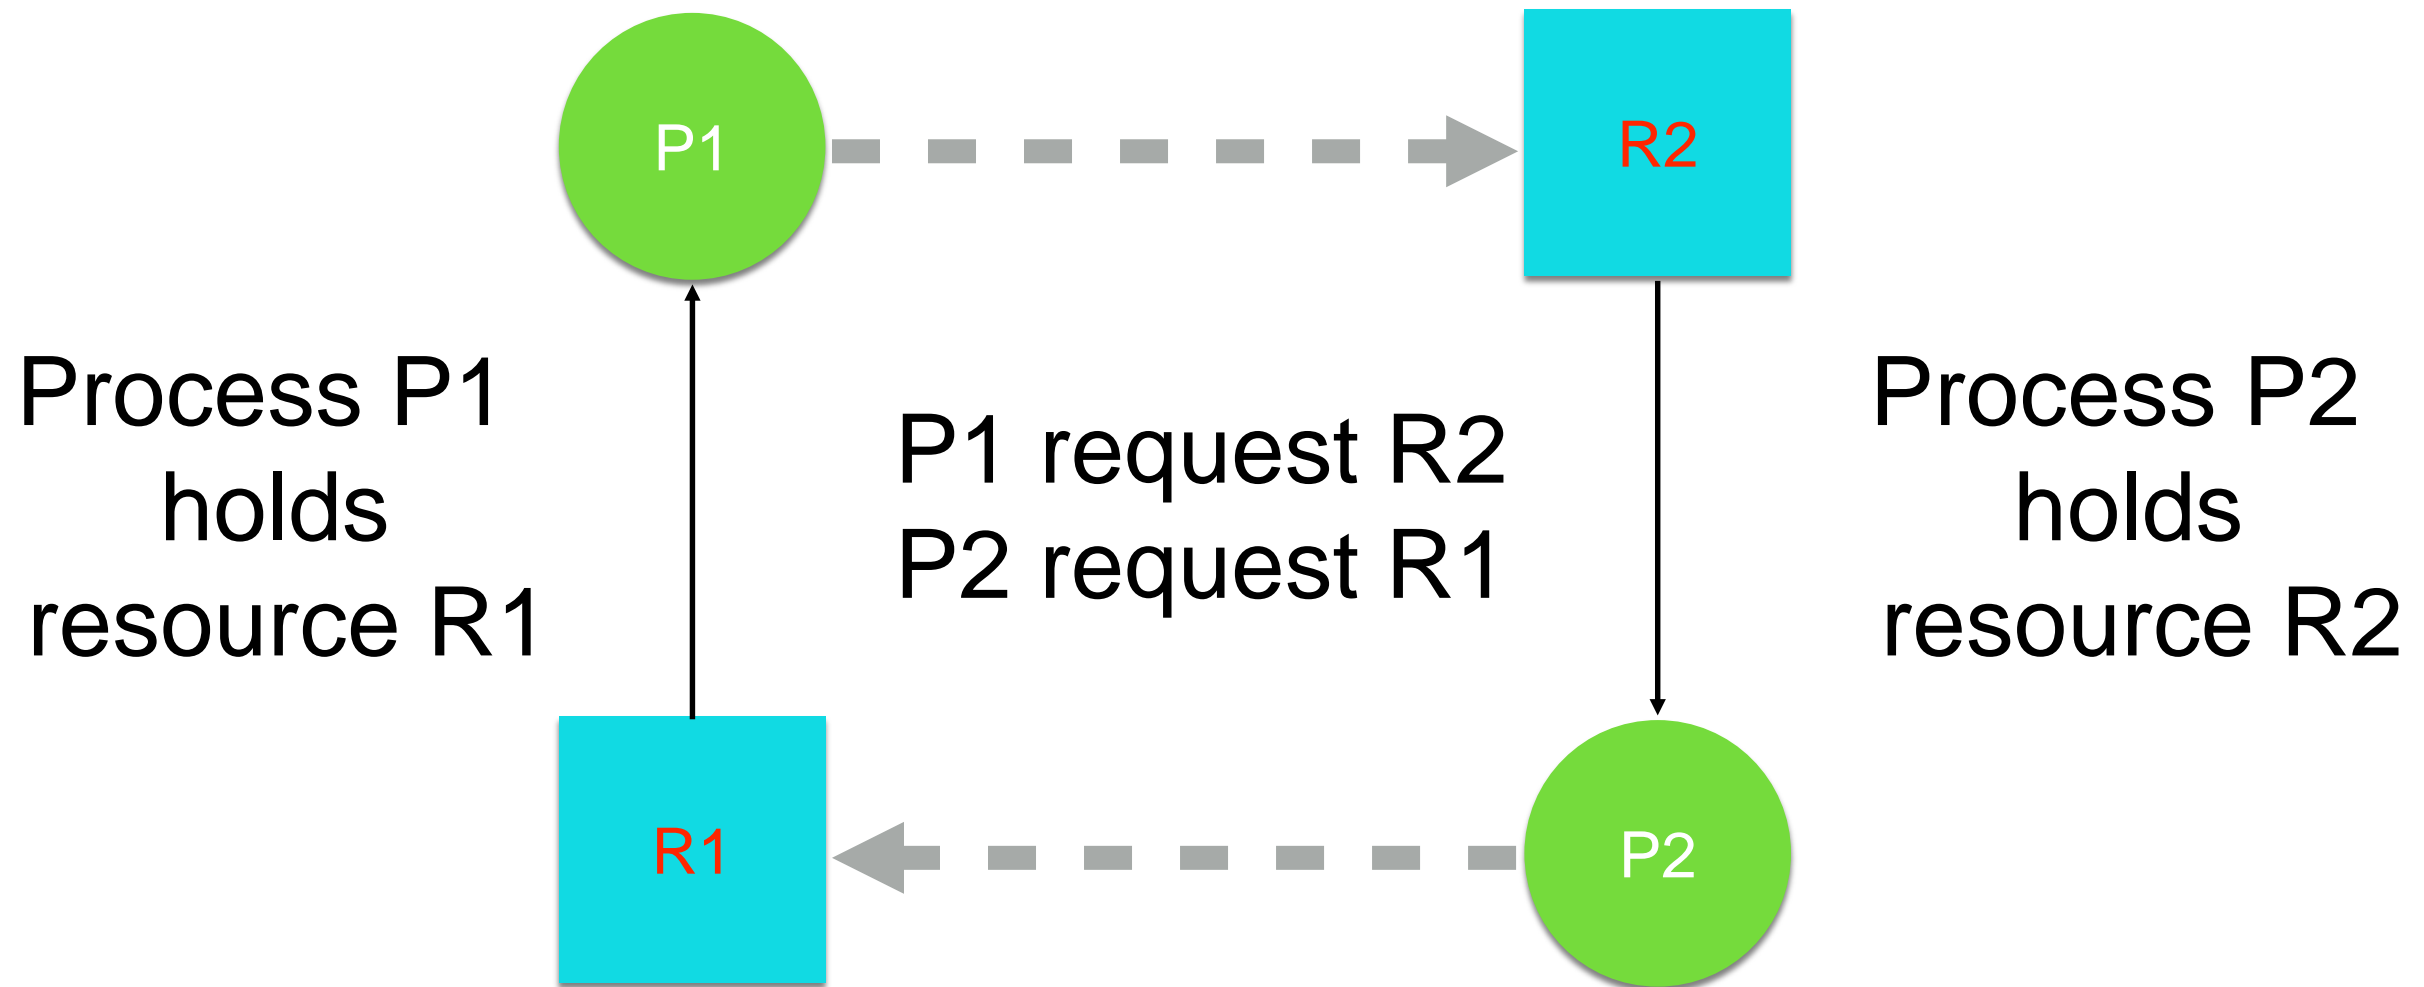


Definition

- **A set of processes is deadlocked if each process in the set is waiting for an event that only another process “ in the set can cause.**

“Modern Operating Systems”, 4rd ed
p.439

Example



Resource allocation graph
if cycle ==> Deadlock

Resources

- All operating systems manage resources
 - CPU, memory, disk, external hardware, ...
- Processes (or threads) sometimes need exclusive access to resources
 - Devices, files, database tables
 - Shared state (counters, queues, lists, etc)
- Preemptable vs nonpreemptable resources
 - Can the resource safely be taken away from a process?
 - Deadlocks typically involve nonpreemptable resources
 - Request ==> Use ==> Release

Types of deadlock

- Resource deadlocks
 - Two processes block waiting for a resource held by the other
- Communication deadlocks
 - Processes block waiting for replies that might never be forthcoming due to dropped packets
- Livelock
 - Processes are running, but not making progress

Deadlock conditions

- Mutual exclusion
 - Only one process may use a resource at any given time
- Hold and wait
 - Process holds at least one resource and waits for a different resource to become available
- No preemption
 - A resource must be released by the process that acquired it
- Circular wait
 - Every process waits for a resource held by the next process in chain




Image credit: Randall Munroe, [http:// xkcd.com/ 506/](http://xkcd.com/506/)

Deadlocks in practice

```
// gcc -o deadlock-threads deadlock-threads.c -lpthread
```

```
#include <pthread.h>
#include <unistd.h>
#include <stdio.h>
#include <stdlib.h>
```

```
pthread_mutex_t  mutex1, mutex2;
int              thread1_counter  = 0,
                thread2_counter  = 0,
                sleep_max_usec    = 1;
```

```
void* thread1(void *arg) {
    while (1) {
        pthread_mutex_lock(&mutex1);
        pthread_mutex_lock(&mutex2);
        thread1_counter++;
        printf("Thread 1: %d\n", thread1_counter);
        pthread_mutex_unlock(&mutex2);
        pthread_mutex_unlock(&mutex1);
        usleep((rand()&0x7fffffff) % sleep_max_usec);
    }
    return 0;
}
```

```
void* thread2(void *arg) {
    while (1) {
        pthread_mutex_lock(&mutex2);
        pthread_mutex_lock(&mutex1);
        thread2_counter++;
        printf("Thread 2: %d\n", thread2_counter);
        pthread_mutex_unlock(&mutex1);
        pthread_mutex_unlock(&mutex2);
        usleep((rand()&0x7fffffff) % sleep_max_usec);
    }
    return 0;
}
```

```
int main(int argc, char *argv[]) {
    pthread_t tid;
    srand(time(0));
    sleep_max_usec = (argc > 1 ? atoi(argv[1]) : 1000);
    pthread_mutex_init(&mutex1, 0);
    pthread_mutex_init(&mutex2, 0);
    pthread_create(&tid, 0, thread1, 0);
    pthread_detach(tid);
    thread2(0);
    return 0;
}
```

Could OS detect this deadlock?
Deadlock to livelock with spinlocks



Handling deadlock

- Detect and recover
 - More common in database systems, not so much in operating systems
- Avoid deadlocks using “smart” resource allocation
- Prevent by negating one of the deadlock conditions
- Ignore the problem (“Ostrich algorithm”)
 - In practice, this is what most operating systems do
 - Why pay a (potentially severe) performance penalty for an issue that occurs infrequently?



```
#!/usr/bin/env python
```

```
graph = ["R", "A", "B", "C", "S", "D",
         "T", "E", "F", "U", "V", "W", "G"]
edges = { "R" : [["A", False]], "A" : [["S", False]],
         "B" : [["T", False]], "C" : [["S", False]],
         "D" : [["S", False], ["T", False]],
         "E" : [["V", False]], "F" : [["S", False]],
         "G" : [["U", False]], "S" : [],
         "T" : [["E", False]], "U" : [["D", False]],
         "V" : [["G", False]], "W" : [["F", False]] }
```

```
def unmarkEdges():
    for node in edges.keys():
        for i in range(0, len(edges[node])):
            edges[node][i][1] = False

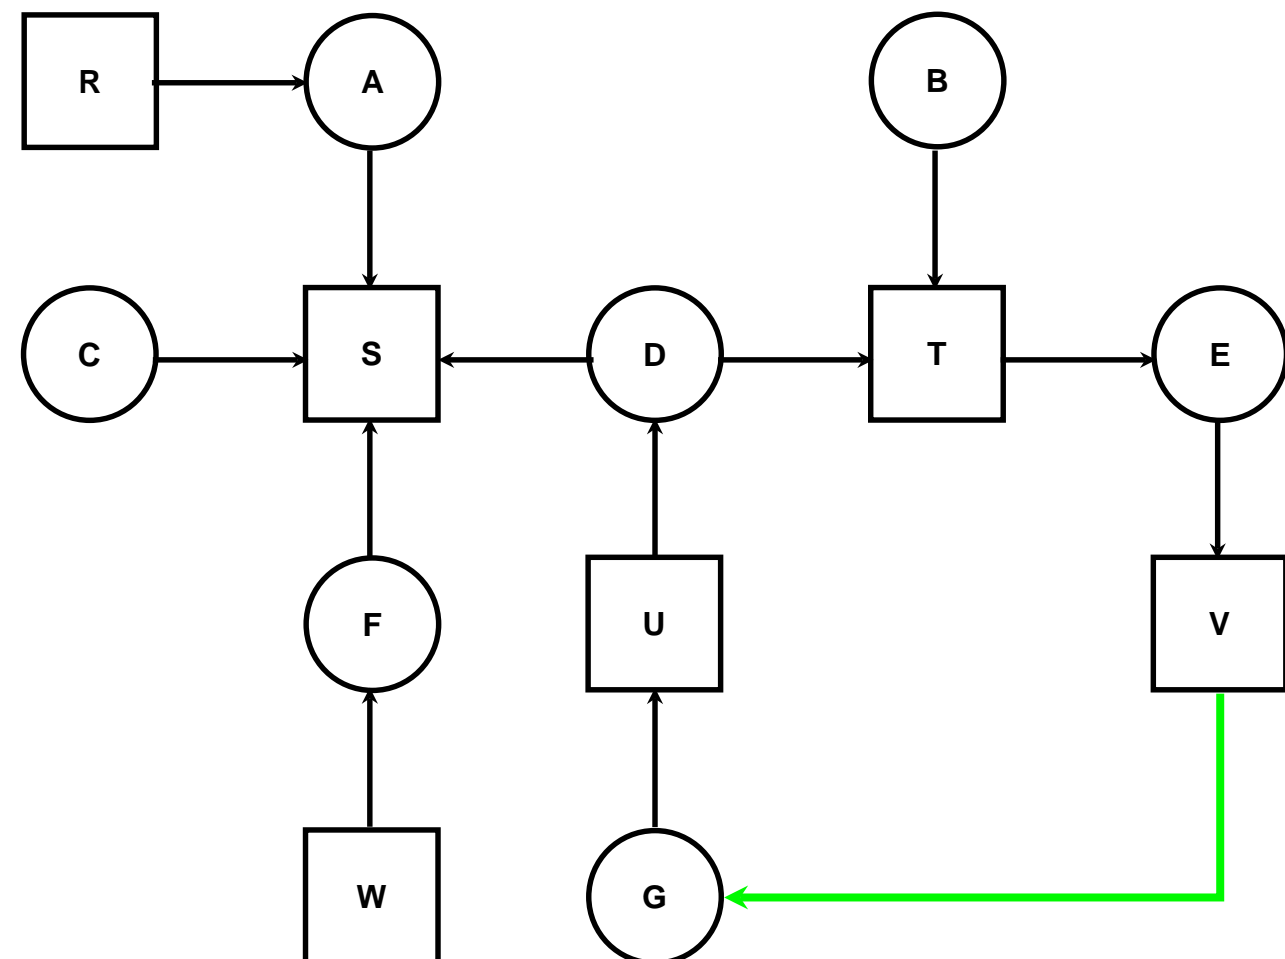
def detectDeadlock():
    for node in graph:
        L = [node]
        unmarkEdges()
        while len(L) > 0:
            currentNode = L[-1]
            foundUnmarked = False
            print "Current: ", currentNode, "List", L
            for edge in edges[currentNode]:
                if not edge[1]:
                    if edge[0] in L:
                        L.append(edge[0])
                        print "Deadlock detected:", L
                        return True
                    L.append(edge[0])
                    edge[1] = True
                    foundUnmarked = True
                    break
            if not foundUnmarked:
                print "No cycle detected starting at ", L[0], ":", L
                L = L[:-1]

    return False

if __name__ == "__main__":
    detectDeadlock()
```

Detection

- Inspect resource allocation graph
- Look for cycles
- Depth-first traversal





Recovery

(hard)

- Preemption
 - Temporarily take a resource away from the process holding it
 - Not usually feasible, depends strongly on characteristics of resource being preempted
- Rollback
 - Processes are checkpointed periodically
 - Deadlock detected? Roll process state back to state before acquisition of a needed resource
- Kill processes
 - Crude example: Linux Out-of-Memory (OOM) killer

Avoidance

- Resources typically not acquired all at once
- Can resource requests be scheduled so as to always avoid deadlock?
 - Yes, if max resource need is known in advance for all processes

Safe and unsafe states

- Assume max number of resources a process will use is known
- Track state of each process' resource usage
- Safe state: A sequence of resource allocations exists that allow all processes to complete
- Unsafe state: No guarantee that all processes will complete without deadlocking

Total: 8

Safe state:

	Has	Max
P ₁	2	6
P ₂	2	3
P ₃	3	5

Free: 1

	Has	Max
P ₁	2	6
P ₂	3	3
P ₃	3	5

Free: 0

	Has	Max
P ₁	2	6
P ₂	0	0
P ₃	3	5

Free: 3

	Has	Max
P ₁	2	6
P ₂	0	0
P ₃	5	5

Free: 1

	Has	Max
P ₁	2	6
P ₂	0	0
P ₃	0	0

Free: 6

Unsafe state:

	Has	Max
P ₁	4	6
P ₂	1	3
P ₃	2	5

Free: 1

?

Either process
may require 5
resources, and
end up
deadlocked

Banker's algorithm

- Only grant resource requests that lead to a new safe state
- Otherwise, delay requests from processes
- Practical utility of algorithm is questionable
 - Max resource count rarely known in advance
 - Resources may disappear
 - New processes are added to the mix

Prevention

- Four conditions required for deadlock
 - (1) Mutual exclusion, (2) hold and wait, (3) no preemption and (4) circular wait
- Mutual exclusion
 - Deadlocks on a printer can be avoided by spooling output; only the printer daemon talks to the printer
 - Deadlocks still possible if printer daemon waits for complete file and all disk space is consumed
 - This time, we deadlock on the disk

Prevention: Hold and wait

- Request all resources needed upfront
 - What if we don't know which resources are needed?
 - Inefficient, as resources are tied up until the process exits
- Release held resources, then attempt to acquire all resources at once

Prevention: No preemption

- Virtualized resources (printer daemon)
- Not really possible for database tables
 - Data corruption/loss

Prevention: Circular wait

- Only one resource at a time per process
 - Need a different resource? Release the other resource first
- Global resource numbering
 - Always acquire resources in ascending order
 - Cycles can never occur
 - Is it practical to assign a strict order of all resources governed by an operating system?



Starvation

- Similar to deadlock/livelock
- Scheduling policy affects possibility of starvation
- For instance, shortest job first may result in longer-running jobs never getting a chance to execute/acquire resources
 - FIFO allocation of resources avoids starvation



Email

Password

Log in



[Forgot your password?](#)

Weather services from Norwegian Meteorological Institute, for governmental agencies

Practical Starvation

- Extreme Weather event «NINA» January 2015
- yr.no hits 6 mill distinctive users and handles this
- halo.met.no is resource for governmental use
 - Server crashes
 - «ulimit» sets maximum number of open files for a process. Set to 10 000. Previously observed around 1 000.
 - Demand requests around 12 000 files.
 - Smart programmer had good error reporting

Debugging deadlocks

- `pthread_setname_np(...)`, `pthread_getname_np(...)`
 - “np” = not portable
 - Debuggers typically display thread name
 - `getname()` can be used for `printf`'s
- `dtrace` and similar tools



Summary

- Deadlock: Set of processes where each is waiting for resources that can only be released by another in same set.
- Four conditions required for deadlock
 - (1) Mutual exclusion, (2) hold and wait, (3) no preemption and (4) circular wait
- Handling:
 - Reboot, remove device, kill application and restart. Checkpoint the application and kill.

Other presentations

- Princeton – Singh
 - <http://www.cs.princeton.edu/courses/archive/fall18/cos318/lectures/9.Deadlock.pdf>

Questions?